

Homework 8: Cache-Oblivious Algorithms

Then, answer the writeup questions in this handout and submit an *individual* writeup. See the following paper for more information on cache-oblivious algorithms:
<https://dl.acm.org/citation.cfm?id=2071383>.

For this homework, assume that all matrices are stored in row-major layout.

1 Cache Complexity of Matrix Multiplication

During Lecture 14 we discussed the cache complexity of matrix multiplication of dimension n , with tall cache assumption of size \mathcal{M} and cache line size \mathcal{B} . For the naive approach, there were two cases: 1) If $n > \mathcal{M}/\mathcal{B}$, then $\Theta(n^3)$ cache misses occur, and 2) if $\mathcal{M}^{1/2} < n < \mathcal{M}/\mathcal{B}$, then $\Theta(n^3/\mathcal{B})$ cache misses occur. For the blocking approach, with block size $s < \mathcal{M}^{1/2}$, $\Theta(n^3/\mathcal{B}\mathcal{M}^{1/2})$ cache misses occur. The cache-oblivious approach achieves the same complexity as the blocking approach without the need of the voodoo parameter s .

Checkoff Item 1: Assume we want to multiply two rectangular matrices: $m \times n$ with $n \times r$. Given the same tall cache assumption, please analyze the complexity for one of the following four cases: the two cases for the naive approach ($n > \mathcal{M}/\mathcal{B}$ and $\mathcal{M}/r < n < \mathcal{M}/\mathcal{B}$), the block approach, and the cache-oblivious approach. You may pick whichever case you want to analyze.

2 Tableau Construction

Consider the tableau-construction problem from Lecture 8. The problem involves filling an $N \times N$ tableau, where each entry of the tableau is calculated as a function of some of its neighbors. To be specific, the equation to fill an element of the tableau would take the form

$$A[i][j] = f(A[i-1][j-1], A[i][j-1], A[i-1][j])$$

where f is an arbitrary function.

2.1 Iterative Formulation

Consider the code snippet in Figure 1 below.

```

01 #define A(i, j) A[N + (i) - (j) - 1]
02
03 void tableau(double *A, size_t N) {
04     for (size_t i = 1; i < N; i++) {
05         for (size_t j = 1; j < N; j++) {
06             A(i, j) = f(A(i-1, j-1), A(i, j-1), A(i-1, j));
07         }
08     }
09 }

```

Figure 1: A simple, iterative loop for filling a tableau.

In this problem, we are only interested in computing the final value of the tableau, stored in $A(N-1, N-1)$, and hence we really only need $2N - 1$ amount of space during computation. Thus, the algorithm declares A as an array of size $2N - 1$.

The algorithm initializes the first row and first column of the tableau, and invokes the `tableau` function as shown in Figure 2.

```

10 for (size_t i = 0; i < N; i++) {
11     A(i, 0) = INIT_VAL;
12 }
13 for (size_t j = 0; j < N; j++) {
14     A(0, j) = INIT_VAL;
15 }
16 tableau(A, N);
17 res = A(N - 1, N - 1);

```

Figure 2: Initializing and calling the iterative `tableau` function.

Write-up 1: Explain why $2N - 1$ space is sufficient and how the `tableau` function utilizes the $2N - 1$ space.

Recall the tall cache assumption, which states that $B^2 < \alpha \mathcal{M}$, where B is the size of the cache line, \mathcal{M} is the size of the cache, and $\alpha \leq 1$ is a constant.

Write-up 2: Assuming that an optimal replacement strategy holds and that the cache is tall, give a tight upper bound on the cache complexity $Q(n)$ for each of the following cases using O notation, where $c \leq 1$ is a sufficiently small constant:

1. $n \geq cM$
2. $n < cM$

2.2 Recursive Formulation

Now consider the code snippet for a recursive tableau implementation, as shown in Figure 3. This

```

18 #define A(i, j) A[N + (i) - (j) - 1]
19
20 void recursive_tableau(double *A, size_t rbegin, size_t rend, size_t cbegin,
21                        size_t cend) {
22     if (rend-rbegin == 1 && cend-cbegin == 1) {
23         size_t i = rbegin, j = cbegin;
24         A(i, j) = f(A(i-1, j-1), A(i, j-1), A(i-1, j));
25     } else {
26         size_t rmid = rend-rbegin > 1 ? (rbegin + (rend-rbegin) / 2) : rend;
27         size_t cmid = cend-cbegin > 1 ? (cbegin + (cend-cbegin) / 2) : cend;
28         recursive_tableau(A, rbegin, rmid, cbegin, cmid);
29         if (cend > cmid)
30             recursive_tableau(A, rbegin, rmid, cmid, cend);
31         if (rend > rmid)
32             recursive_tableau(A, rmid, rend, cbegin, cmid);
33         if (rend > rmid && cend > cmid)
34             recursive_tableau(A, rmid, rend, cmid, cend);
35     }
36 }

```

Figure 3: A recursive implementation for filling in a tableau.

algorithm similarly uses only $2N - 1$ amount of space, initializes the array A , and invokes the `recursive_tableau` function as shown in Figure 4. This recursive algorithm divides the tableau into four quadrants to compute. As discussed in Lecture 8 (slide 88), after the first quadrant is done computing, we can then compute the second and third quadrants in parallel. Parallelizing this way gives us work as $\Theta(n^2)$ and span as $\Theta(n^{\lg 3})$ with parallelism as $\Theta(n^{2-\lg 3})$. We also discussed (slide 92) a more parallel construction that divides up the tableau 9 ways.

```

37 for (size_t i = 0; i < N; i++) {
38     A(i, 0) = INIT_VAL;
39 }
40 for (size_t j = 0; j < N; j++) {
41     A(0, j) = INIT_VAL;
42 }
43 if (N > 1) {
44     recursive_tableau(A, 1, N, 1, N);
45 }
46 res = A(N-1, N-1);

```

Figure 4: Initializing and calling the `recursive_tableau` function.

Write-up 3: Derive the general formula for work and span, assuming a k^2 -way tableau construction (i.e., the tableau is divided up into k^2 pieces of size $n/k \times n/k$).

Write-up 4: Answer the following questions assuming that an optimal replacement strategy holds and that the cache is tall.

1. Show the recurrence relation for the cache complexity $Q(n)$ using the 4-way construction of the `recursive_tableau` function.
2. Draw the recursion tree and label the internal nodes and leaves with their cache complexity $Q(n)$. What's the height of the recursion tree?
3. How many leaves are in the recursion tree?
4. Using the recursion tree and the recurrence relation, derive a simplified expression for $Q(n)$.

Write-up 5: Answer the following question assuming that an optimal replacement strategy holds and that the cache is tall. Assuming a k^2 -way tableau construction, show that if we are “unlucky,” where a subpiece is just slightly above the cache size, then we have $Q(n) = \Theta(n^2 k / \mathcal{MB})$. Also show that if we are lucky and this situation does not arise, then we have $Q(n) = \Theta(n^2 / \mathcal{MB})$.