


6.172
Performance
Engineering of
Software
Systems



LECTURE 18
Domain Specific Languages and
Autotuning

Saman Amarasinghe

Domain Specific Languages


- Capture the programmer intent at a higher level of abstraction
- Obtain many software engineering benefits
 - clarity, portability, maintainability, testability, etc...
- Provide the compiler more opportunities for higher performance
 - Can encode expert knowledge of domain specific transformations
 - Better view of the computation performed without heroic analysis
 - Less low-level decisions by the programmer that has to be undone

Outline

- GraphIt
- Halide
- OpenTuner


PageRank Example in C++


```
void pagerank(Graph &graph, double * new_rank, double * old_rank, int * out_degree, int max_iter){  
    for (i = 0; i < max_iter; i++) {  
        for (src : graph.vertices()) {  
            for (dst : graph.getOutgoingNeighbors(node)) {  
                new_rank[dst] += old_rank[src]/out_degree[src]; } } }  
        for (node : graph.vertices()) {  
            new_rank[node] = base_score + damping*new_rank[node]; } }  
    swap (old_rank, new_rank); }  
}
```



PageRank Example in C++


```
void pagerank(Graph &graph, double * new_rank, double * old_rank, int * out_degree, int max_iter){  
    for (i = 0; i < max_iter; i++) {  
        for (src : graph.vertices()) {  
            for (dst : graph.getOutgoingNeighbors(node)) {  
                new_rank[dst] += old_rank[src]/out_degree[src]; } } }  
    for (node : graph.vertices()) {  
        new_rank[node] = base_score + damping*new_rank[node]; }  
    swap (old_rank, new_rank); }  
}
```





PageRank Example in C++

```
void pagerank(Graph &graph, double * new_rank, double * old_rank, int * out_degree, int max_iter){  
    for (i = 0; i < max_iter; i++) {  
        for (src : graph.vertices()) {  
            for (dst : graph.getOutgoingNeighbors(node)) {  
                new_rank[dst] += old_rank[src]/out_degree[src]; } } }  
    for (node : graph.vertices()) {  
        new_rank[node] = base_score + damping*new_rank[node]; } }  
    swap (old_rank, new_rank); }  
}
```



Hand-Optimized C++

```
template<typename APPLY_FUNC>
void edgeset_apply_pull_parallel(Graph &g, APPLY_FUNC apply_func) {
    int64_t numVertices = g.num_nodes(), numEdges = g.num_edges();
    parallel_for(int n = 0; n < numVertices; n++) {
        for (int socketId = 0; socketId < omp_get_num_places(); socketId++) {
            local_new_rank[socketId][n] = new_rank[n]; } }
    int numPlaces = omp_get_num_places();
    int numSegments = g.getNumSegments("s1");
    int segmentsPerSocket = (numSegments + numPlaces - 1) / numPlaces;
    #pragma omp parallel num_threads(numPlaces) proc_bind(spread){
        int socketId = omp_get_place_num();
        for (int i = 0; i < segmentsPerSocket; i++) {
            int segmentId = socketId + i * numPlaces;
            if (segmentId >= numSegments) break;
            auto sg = g.getSegmentedGraph(std::string("s1"), segmentId);
            #pragma omp parallel num_threads(omp_get_place_num_procs(socketId)) proc_bind(close){
                #pragma omp for schedule(dynamic, 1024)
                for (NodeID localId = 0; localId < sg->numVertices; localId++) {
                    NodeID d = sg->graphId[localId];
                    for (int64_t ngh = sg->vertexArray[localId]; ngh < sg->vertexArray[localId + 1]; ngh++) {
                        NodeID s = sg->edgeArray[ngh];
                        local_new_rank[socketId][d] += contrib[s]; }}}
        parallel_for(int n = 0; n < numVertices; n++) {
            for (int socketId = 0; socketId < omp_get_num_places(); socketId++) {
                new_rank[n] += local_new_rank[socketId][n]; }}}
    struct updateVertex {
        void operator()(NodeID v) {
            double old_score = old_rank[v];
            new_rank[v] = (beta_score + (damp * new_rank[v]));
            error[v] = fabs((new_rank[v] - old_rank[v]));
            old_rank[v] = new_rank[v];
            new_rank[v] = ((float) 0); }; };
    void pagerank(Graph &g, double *new_rank, double *old_rank, int *out_degree, int max_iter) {
        for (int i = (0); i < (max_iter); i++) {
            parallel_for(int v_iter = 0; v_iter < builtin_getVertices(edges); v_iter++) {
                contrib[v] = (old_rank[v] / out_degree[v]);};
            edgeset_apply_pull_parallel(edges, updateEdge());
            parallel_for(int v_iter = 0; v_iter < builtin_getVertices(edges); v_iter++) {
                updateVertex()(v_iter); }; }}
```

More than 23x faster
Intel Xeon E5-2695 v3 CPUs with 12 cores
each for a total of 24 cores.

Multi-Threaded


Load Balanced

NUMA Optimized

Cache Optimized


(1) Hard to write correctly.
(2) Extremely difficult to experiment
with different combinations of
optimizations

Graph Algorithms




- What are the different types of graph algorithms

Topology-Driven Algorithms




Work on All Edges and Vertices


Data-Driven Algorithms



Data-Driven Algorithms




Data-Driven Algorithms




**Work on a subset of vertices and edges
(Data-Driven)**

Data-Driven Algorithms




Work on a subset of vertices and edges
(Data-Driven)

Graph Traversal



- Different Traversal Orders have different performance characteristics


Push Traversal



Incurs overhead with atomics

Traverses no extra edges


Pull Traversal



Incurs no overhead from atomics

Traverses extra edges


Partitioning




Improves locality

Needs extra instructions to traverse two graphs

Power-Law Graphs



World Wide Web




Maps

© [Google](#). All rights reserved. This content is excluded from our Creative Commons license.
For more information, see <https://ocw.mit.edu/help/faq-fair-use/>

**Power-Law Degree Distribution,
Small Diameter, Poor Locality**


Social Networks




Engineering Meshes

This image is in the [public domain](#).

Bounded-Degree Graphs




World Wide Web




Maps

© Google. All rights reserved. This content is excluded from our Creative Commons license.
For more information, see <https://ocw.mit.edu/help/faq-fair-use/>

Power-Law Degree Distribution,
Small Diameter, Poor Locality




Bounded Degree Distribution
Large Diameter, Excellent Locality




Engineering Meshes


This image is in the [public domain](#).

Optimization Tradeoff Space







Graphs




Algorithms




Hardware

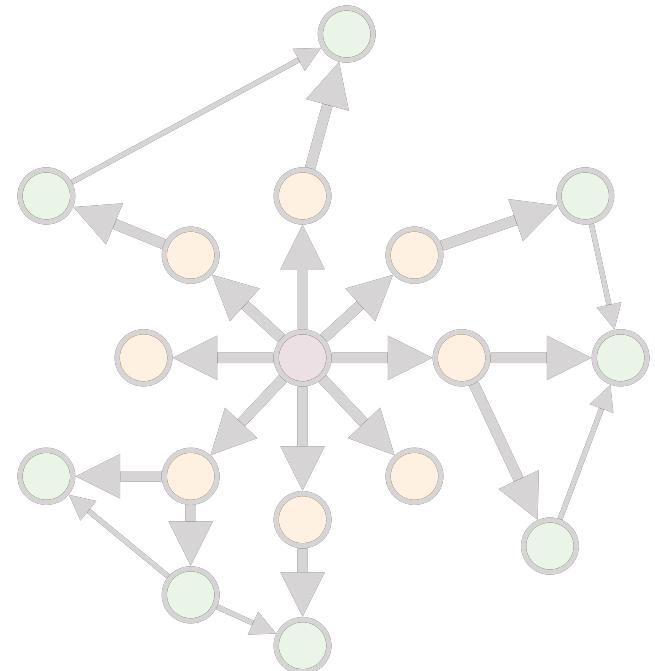
Optimizations



Graphs




Algorithms



A large orange triangle is inscribed within a black circle. Three small colored dots are placed on the triangle's surface: a blue dot near the top vertex, a red dot near the middle left side, and a green dot near the bottom left vertex.


Pull Partitioning Vertex-Parallel

Hardware



Graphs


Algorithms




A large orange triangle is inscribed within a black circle. The triangle's vertices are in contact with the circle's circumference. A small green dot is positioned on the left side of the triangle, roughly one-third of the way up from the bottom vertex. A small black dot is located near the top center of the triangle, slightly to the right of the vertical midpoint.

Hardware


Push Vertex-Parallel




Graphs



Algorithms



**Bad
optimizations
(schedules) can
be > 100x slower**



Optimizations

Hardware

**Pull
Partitioning
Vertex-Parallel**

GraphIt

A Domain-Specific Language for Graph Applications

- **Decouple algorithm from optimization for graph applications**


- **Algorithm:** What to Compute

- **High level** ignores all the optimization details


- **Optimization (schedule):** How to Compute

- **Easy to use** for users to try different combinations


- **Powerful enough** to beat hand-optimized libraries by up to 4.8x




Algorithm Language



`edges.apply(func)`





`edges.from(vertexset)
.to(vertexset)
.srcFilter(func)
.dstFilter(func)
.apply(func)`



`vertices.apply(func)`

PageRank Example





```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

PageRank Example





```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

PageRank Example



```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end
```


```
func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end
```

```
func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

GraphIt

A Domain-Specific Language for Graph Applications

- **Decouple algorithm from optimization for graph applications**
- **Algorithm:** What to Compute
 - **High level** ignores all the optimization details



Scheduling Language

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Scheduling Language

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Schedule 1

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Scheduling Functions

schedule:

```
program->configApplyDirection("s1", "SparsePush");
```

Schedule 1

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Scheduling Functions

```
schedule:
    program->configApplyDirection("s1", "SparsePush");
```

Schedule 1

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Pseudo Generated Code

```
double * new_rank = new double[num_verts];
double * old_rank = new double[num_verts];
int * out_degree = new int[num_verts];

...
for (NodID src : vertices) {
    for(NodID dst : G.getOutNgh(src)){
        new_rank[dst] += old_rank[src] / out_degree[src];
    }
}
....
```

Scheduling Functions

```
schedule:
    program->configApplyDirection("s1", "SparsePush");
```

Schedule 2

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Pseudo Generated Code

```
double * new_rank = new double[num_verts];
double * old_rank = new double[num_verts];
int * out_degree = new int[num_verts];

...
parallel_for (NodeID src : vertices) {
    for(NodeID dst : G.getOutNgh(src)){
        atomic_add (new_rank[dst],
                    old_rank[src] / out_degree[src] );
    }
}
....
```

Scheduling Functions

schedule:

```
program->configApplyDirection("s1", "SparsePush");
program->configApplyParallelization("s1", "dynamic-vertex-parallel");
```

Schedule 3

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Pseudo Generated Code

```
double * new_rank = new double[num_verts];
double * old_rank = new double[num_verts];
int * out_degree = new int[num_verts];

...
parallel_for (NodeID dst : vertices) {
    for(NodeID src : G.getInNgh(dst)){
        new_rank[dst] += old_rank[src] / out_degree[src];
    }
}
....
```

Scheduling Functions

```
schedule:
program->configApplyDirection("s1", "DensePull");
program->configApplyParallelization("s1", "dynamic-vertex-parallel");
```

Schedule 4

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:max_iter
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Pseudo Generated Code


```
double * new_rank = new double[num_verts];
double * old_rank = new double[num_verts];
int * out_degree = new int[num_verts];

...
for (Subgraph sg : G.subgraphs) {
    parallel_for (NodeID dst : verticesa) {
        for(NodeID src : G.getInNgh(dst)){
            new_rank[dst] += old_rank[src] / out_degree[src];
        }
    }
}
....
```

Scheduling Functions

```
schedule:
program->configApplyDirection("s1", "DensePull");
program->configApplyParallelization("s1", "dynamic-vertex-parallel");
program->configApplyNumSSG("s1", "fixed-vertex-count", 10);
```

Speedups of Schedules



Intel Xeon E5-2695 v3 CPUs with 12 cores each for a total of 24 cores and 48 hyper-threads.

Many More Optimizations

- Direction optimizations (`configApplyDirection`),
 - **SparsePush, DensePush, DensePull, DensePull-SparsePush, DensePush-SparsePush**
- Parallelization strategies (`configApplyParallelization`)
 - **serial, dynamic-vertex-parallel, static-vertexparallel, edge-aware-dynamic-vertex-parallel, edge-parallel**
- Cache (`configApplyNumSSG`)
 - **fixed-vertex-count, edge-aware-vertexcount**
- NUMA (`configApplyNUMA`)
 - **serial, static-parallel, dynamic-parallel**
- AoS, SoA (`fuseFields`)
- Vertexset data layout (`configApplyDenseVertexSet`)
 - **bitvector, boolean array**

State of the Art and GraphIt

	FT	RD	WB	TW	LJ
	PR	BFS	CC	SSSP	
Ligra (PPoPP13)					
FT	3.48	1	1	1	
RD	5.63	1.13	3.12	1.14	
WB	4.15	1.42	2.96	1.13	
TW	2.69	4.81	2.16	4.57	
LJ	6.17	1.38	4.94	2.77	

	FT	RD	WB	TW	LJ
	PR	BFS	CC	SSSP	
GraphMat (VLDB15)					
FT	1.64	3.7	5.98	1.86	
RD	2.34	9.4	11	1.62	
WB	2.14	7.44	9.13	2.98	
TW	1.61	9.06	7.04	151	
LJ					

	FT	RD	WB	TW	LJ
	PR	BFS	CC	SSSP	
GreenMarl (ASPLOS12)					
FT	1.51	1.83	3.06	1.82	
RD	2.42	6.03	5.78	1.41	
WB	2.59	2.84	5.96	2.54	
TW	1.26	2.45	8.99	328	
LJ					

	FT	RD	WB	TW	LJ
	PR	BFS	CC	SSSP	
GraphIt (OOPSLA18)					
FT	1	1.3	1.11	1.07	
RD	1	1	1	1	
WB	1	1	1	1	
TW	1.23	1	1.43	1	
LJ	1	1	1	1	

Intel Xeon E5-2695 v3 CPUs with 12 cores each for a total of 24 cores and 48 hyper-threads.

Halide

- A new language & compiler
 - Originally developed for image processing
 - Focuses on stencils on regular grids
 - Complex pipelines of stencil kernels
 - Support other operations like reductions and scans
- Primary goal
 - Match or exceed hand optimized performance on each architecture
 - Reduce the rote programming burden of highly optimized code
 - Increase the portability without loss of performance

A Simple Example: 3X3 Blur

```
void box_filter_3x3(const Image &in, Image &blury) {  
  
    Image blurx(in.width(), in.height()); // allocate blurx array  
  
    for (int y = 0; y < in.height(); y++)  
        for (int x = 0; x < in.height(); x++)  
            blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;  
  
    for (int x = 0; x < in.width(); x++)  
        for (int y = 0; y < in.height(); y++)  
            blury(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;
```

Hand-Optimized C++

9.9 → 0.9 ms/megapixel

11x faster
(quad core
x86)

```
void box_filter_3x3(const Image &in, Image &blury) {
    __m128i one_third = _mm_set1_epi16(21846);
    #pragma omp parallel for
    for (int yTile = 0; yTile < in.height(); yTile += 32) {
        __m128i a, b, c, sum, avg;
        __m128i blurx[(256/8)*(32+2)]; // allocate tile blurx array
        for (int xTile = 0; xTile < in.width(); xTile += 256) {
            __m128i *blurxPtr = blurx;
            for (int y = -1; y < 32+1; y++) {
                const uint16_t *inPtr = &(in[yTile+y][xTile]);
                for (int x = 0; x < 256; x += 8) {
                    a = _mm_loadu_si128((__m128i*)(inPtr-1));
                    b = _mm_loadu_si128((__m128i*)(inPtr+1));
                    c = _mm_load_si128((__m128i*)(inPtr));
                    sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                    avg = _mm_mulhi_epi16(sum, one_third);
                    _mm_store_si128(blurxPtr++, avg);
                    inPtr += 8;
                }
                blurxPtr = blurx;
                for (int y = 0; y < 32; y++) {
                    __m128i *outPtr = ((__m128i *)(&(blury[yTile+y][xTile])));
                    for (int x = 0; x < 256; x += 8) {
                        a = _mm_load_si128(blurxPtr+(2*256)/8);
                        b = _mm_load_si128(blurxPtr+256/8);
                        c = _mm_load_si128(blurxPtr++);
                        sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
                        avg = _mm_mulhi_epi16(sum, one_third);
                        _mm_store_si128(outPtr++, avg);
                    }
                }
            }
        }
    }
}
```

Tiled, fused

Vectorized

Multithreaded

Redundant
computation

Near roof-line
optimum

Local Laplacian Filters

prototype for Adobe Photoshop Camera Raw / Lightroom

Reference: 300 lines C++

Adobe: 1500 lines

3 months of work

10x faster (vs. reference)

Halide: 60 lines


1 intern-day

20x faster (vs. reference)

2x faster (vs. Adobe)

GPU: 90x faster (vs. reference)

9x faster (vs. Adobe)



Decouple Algorithm From Schedule

- **Algorithm:** *what* is computed
 - The algorithm defines pipelines as pure functions
 - Pipeline stages are functions from coordinates to values
 - Execution order and storage are unspecified


`blurx(x, y) = (in(x-1, y) + in(x, y) + in(x+1, y))/3;`

`blury(x, y) = (blurx(x, y-1) + blurx(x, y) + blurx(x, y+1))/3;`

Decouple Algorithm From Schedule


- **Algorithm:** *what* is computed
 - The algorithm defines pipelines as pure functions
 - Pipeline stages are functions from coordinates to values
 - Execution order and storage are unspecified
- **Schedule:** *where* and *when* it's computed
 - Architecture Specific
 - Single, unified model for all schedules
 - Simple enough to search, expose to user
 - Powerful enough to beat expert-tuned code

Stencil Pipelines Require Tradeoffs Determined By Organization Of Computation




Parallelism


- Need parallelism to keep multicores, vector units, clusters and GPUs busy
 - Too much parallelism is at best useless but can even be detrimental
- Example: Parallelism of on 3 cores




Too much parallelism



Too little parallelism





Parallelism just right




Locality

- Once a data is touched, how quickly is it reused
- Faster reuse means better cache locality
- Locality at multiple levels: registers, L1, L2, LLC





Too little locality




Good Locality

Redundant Work


- Sometimes cannot get both locality and parallelism
- A little redundant computation can facilitate both
- Extra cost should be amortizable by the wins




Tradeoff Space Modeled By Granularity Of Interleaving




Tradeoff Space Modeled By Granularity Of Interleaving





Tradeoff Space Modeled By Granularity Of Interleaving



Tradeoff Space Modeled By Granularity Of Interleaving




Tradeoff Space Modeled By Granularity Of Interleaving




```
blur_y.tile(xo,yo,xi,yi,W,H)  
blur_x.compute_at(blury, xo)  
.compute_at(blury, xo)
```

Tradeoff Space Modeled By Granularity Of Interleaving



Schedule Primitives Compose To Create Many Organizations



The Bilateral Grid

[Chen et al. 2007]




An accelerated bilateral filter

Original: 122 lines of (clean) C++

Halide: 34 lines of algorithm


On the CPU, 5.9x faster

On the GPU, 2x faster than Chen's hand-written CUDA version



“Snake” Image Segmentation

[Li et al. 2010]



Segments objects in an image using level-sets




Original: 67 lines of MATLAB

Halide: 148 lines of algorithm

On the CPU, 70x faster

MATLAB is memory-bandwidth limited

On the GPU, 1250x faster



Local Laplacian Filters

prototype for Adobe Photoshop Camera Raw / Lightroom

Reference: 300 lines C++

Adobe: 1500 lines

3 months of work

10x faster (vs. reference)

Halide: 60 lines


1 intern-day

20x faster (vs. reference)

2x faster (vs. Adobe)

GPU: 90x faster (vs. reference)


9x faster (vs. Adobe)



Real-World Adoption



Data center
Android
Browser
Glass
G Photos auto-enhance
HDR+
> 200 pipelines, 10s of kLOC in production



Front end language for the
Snapdragon Image Signal Processor



Halide touches every video uploaded.
65B frames/day



OpenTuner

- Performance Engineering is all about finding the right:
 - block size in matrix multiply (voodoo parameters)
 - strategy in the dynamic memory allocation project
 - flags in calling GCC to optimize the program
 - schedule in Halide
 - schedule in GraphIt

How to find the right value

1. Model-Based
2. Heuristic-Based
3. Exhaustive Search
4. Autotuned (OpenTuner)

1. Model Based Solutions

Come-up with a comprehensive model


- In this case, a model for the memory system and data reuse

Pros:

- Can explain exactly why we chose a given tile size
- “Optimal”

Cons:

- Hard to build models
- Cannot model everything
- Our model may miss an important component



2. Heuristic Based Solutions

“A rule of thumb” that works most of the time

- In this case, small two-to-the-power tile sizes works most of the time
- Hard-code them (eg: $S = 8$)

Pros

- Simple and easy to do
- Works most of the time

Cons

- Simplistic
- However, always suboptimal performance
- In some cases may be really bad

3. Exhaustive Search

Empirically evaluate all the possible values


- All possible integers for S

Pros:

- Will find the “optimal” value

Cons:

- Only for the inputs evaluated
- Can take a looooong time!
 - Prune the search space
 - ◆ Only integers that are powers-of-2 from vector register size to the cache size?




4. Autotuning based solutions

What?

1. Define a space of acceptable values
2. Choose a value at random from that space
3. Evaluate the performance given that value
4. If satisfied
 When? → finish
 time limit exceeded → finish
 How?
5. Choose a new value from the feedback
6. Goto 3

Autotuning A Program



Ensembles of techniques

Nelder-Mead
Simplex
(Hill climber)


Differential
Evolution

Particle
Swarm
Optimization

Model Driven
Optimization

- Many different techniques
- Each best suited to solve different problems
- Hard to write a single autotuner that performs well in different domains
- Can we make these techniques work together?

Ensembles of techniques



- Meta-technique divides testing budget between sub-techniques
- Results are shared between all techniques

Autotuning GraphIt

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:11
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Finding the best schedule can
be hard for non-experts.

Scheduling Functions

schedule:

```
program->configApplyDirection("s1", "DensePull");
program->configApplyParallelization("s1", "dynamic-vertex-parallel");
program->configApplyNumSSG("s1", "fixed-vertex-count", 10);
```

Goal

Algorithm Specification

```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:11
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Ideally, the user only need
to write the algorithm

Autotuner

Algorithm Specification


```
func updateEdge (src: Vertex, dst: Vertex)
    new_rank[dst] += old_rank[src] / out_degree[src]
end

func updateVertex (v: Vertex)
    new_rank[v] = beta_score + 0.85*new_rank[v];
    old_rank[v] = new_rank[v];
    new_rank[v] = 0;
end

func main()
    for i in 1:11
        #s1# edges.apply(updateEdge);
        vertices.apply(updateVertex);
    end
end
```

Scheduling Functions


```
schedule:
    program->configApplyDirection("s1", "DensePull");
    program->configApplyParallelization("s1", "dynamic-vertex-parallel");
    program->configApplyNumSSG("s1", "fixed-vertex-count", 10);
```




Input Graphs



Autotuner




Autotuner



Autotuner

**Uses an ensemble
of search methods.
Build on top of
OpenTuner
[PACT14]**



Autotuner

Finds a few schedules that outperform hand-tuned schedules

