

6.172  
Performance  
Engineering  
of Software  
Systems




LECTURE 11  
Storage  
Allocation

Julian Shun




STACKS



# Stack Allocation

## Array and pointer




Allocate  $x$  bytes

```
sp += x;  
return sp - x;
```

# Stack Allocation

## Array and pointer




Allocate  $x$  bytes

```
sp += x;  
return sp - x;
```

Should check for  
stack overflow.

# Stack Allocation

## Array and pointer




Allocate  $x$  bytes

```
sp += x;  
return sp - x;
```

Should check for  
stack overflow.

# Stack Deallocation

## Array and pointer



Allocate  $x$  bytes


```
sp += x;  
return sp - x;
```

Free  $x$  bytes

```
sp -= x;
```

# Stack Deallocation

## Array and pointer



Allocate  $x$  bytes

```
sp += x;  
return sp - x;
```


Free  $x$  bytes

```
sp -= x;
```

Should check for  
stack underflow.

# Stack Storage

## Array and pointer



Allocate  $x$  bytes

```
sp += x;  
return sp - x;
```

Free  $x$  bytes

```
sp -= x;
```

- Allocating and freeing take  $\Theta(1)$  time.
- Must free consistent with stack discipline.
- Limited applicability, but great when it works!
- One can allocate on the call stack using `alloca()`, but this function is deprecated, and the compiler is more efficient with fixed-size frames.



# Stacks and Heaps




Image is in the public domain.


**Stack**



Image is in the public domain.

**Heap**

# FIXED-SIZE HEAP ALLOCATION



# Heap Allocation\*

C provides `malloc()` and `free()`.

C++ provides `new` and `delete`.

Unlike Java and Python, C and C++ provide no **garbage collector**. Heap storage allocated by the programmer must be freed explicitly.

Failure to do so creates a **memory leak**. Also, watch for **dangling pointers** and **double freeing**.


**Memory checkers** (e.g., AddressSanitizer, Valgrind) can assist in finding these pernicious bugs.

---

\*Do not confuse with a **heap data structure**.

# Fixed-Size Allocation

## Free list




- Every piece of storage has the same size
- Unused storage has a pointer to next unused block

## Bitmap mechanism

- Bit for each block saying whether or not it is free
- Bit tricks for allocation

# Fixed-Size Allocation

## Free list




Allocate 1 object

```
x = free;  
free = free->next;  
return x;
```

# Fixed-Size Allocation

## Free list




## Allocate 1 object

```
x = free;  
free = free->next;  
return x;
```

# Fixed-Size Allocation

## Free list



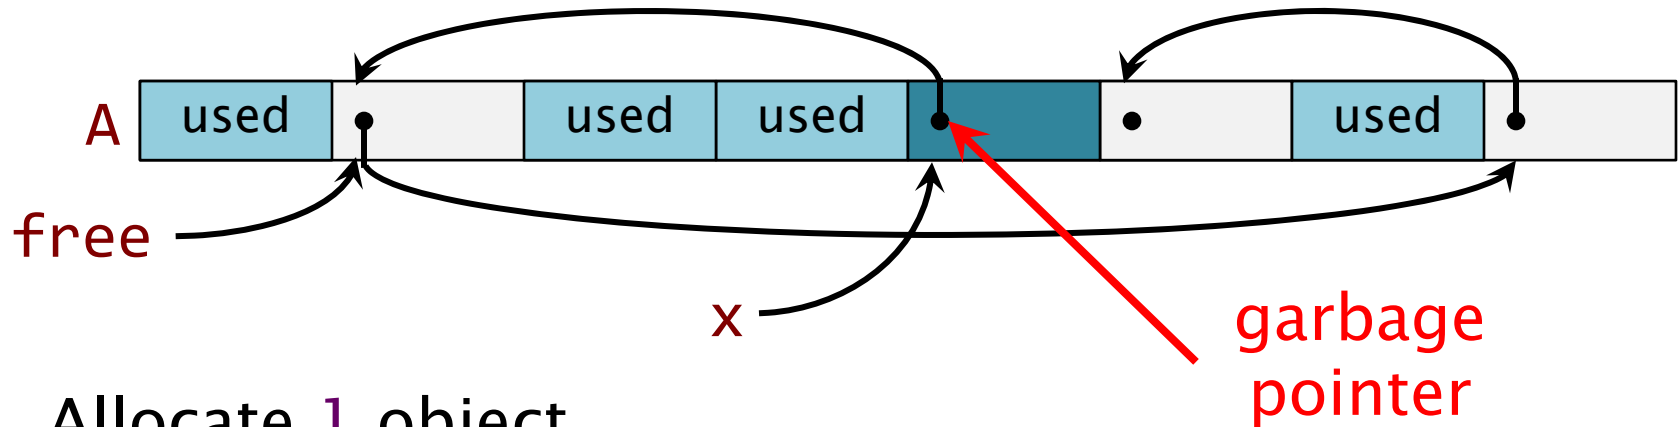
## Allocate 1 object

```
x = free;  
free = free->next;  
return x;
```

Should check  
`free != NULL.`

# Fixed-Size Allocation

## Free list



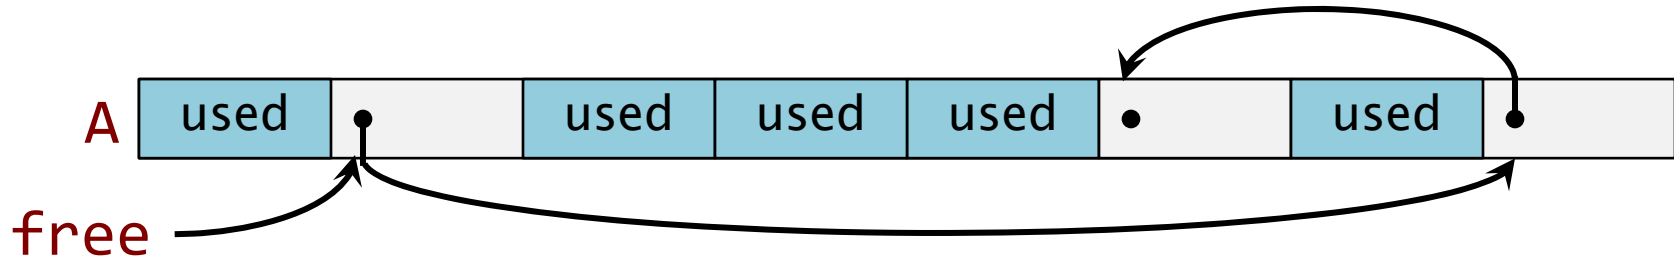
Allocate 1 object

```
x = free;  
free = free->next;  
return x;
```



# Fixed-Size Allocation

## Free list

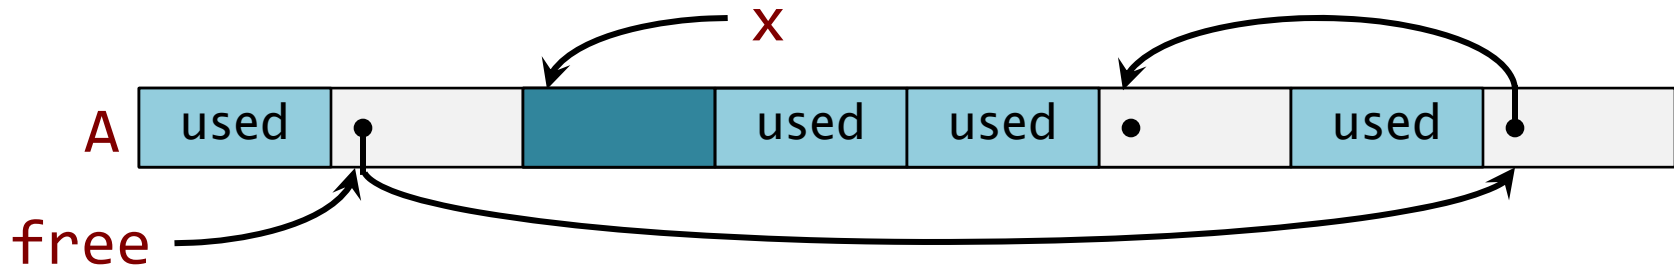


Allocate 1 object

```
x = free;  
free = free->next;  
return x;
```

# Fixed-Size Deallocation

## Free list



## Allocate 1 object

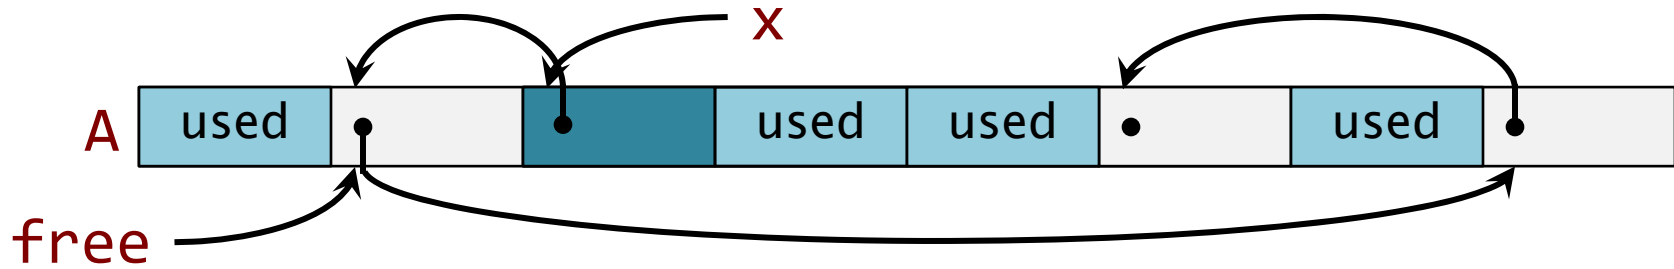
```
x = free;  
free = free->next;  
return x;
```

## free object x

```
x->next = free;  
free = x;
```

# Fixed-Size Deallocation

Free list



Allocate **1** object

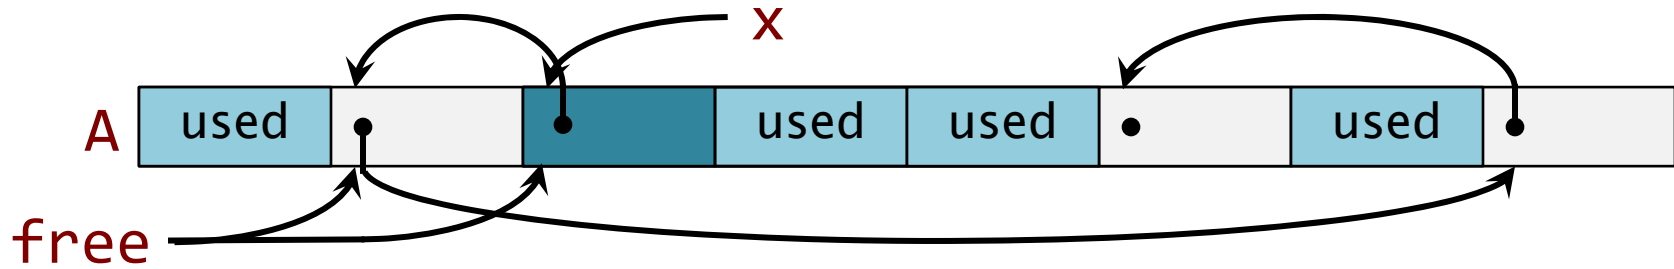
```
x = free;  
free = free->next;  
return x;
```

free object **x**

```
x->next = free;  
free = x;
```

# Fixed-Size Deallocation

## Free list



## Allocate 1 object

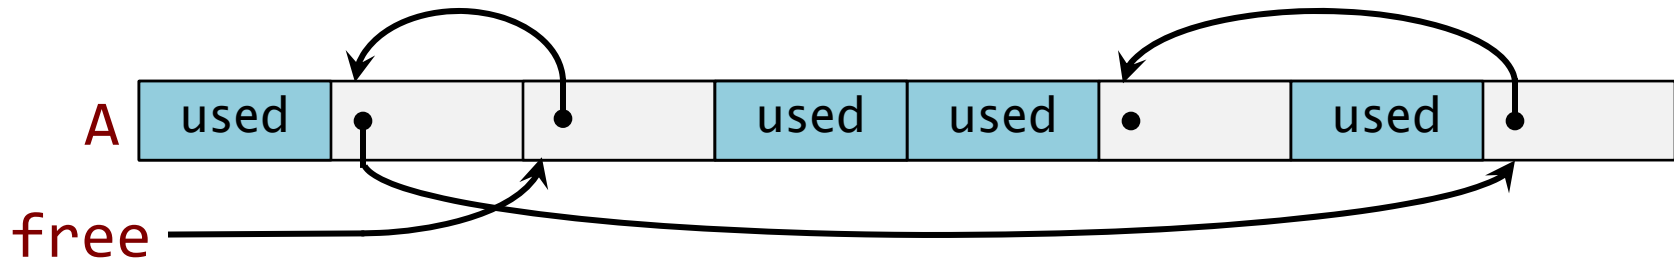
```
x = free;  
free = free->next;  
return x;
```

## free object x

```
x->next = free;  
free = x;
```

# Fixed-Size Deallocation

## Free list



## Allocate 1 object

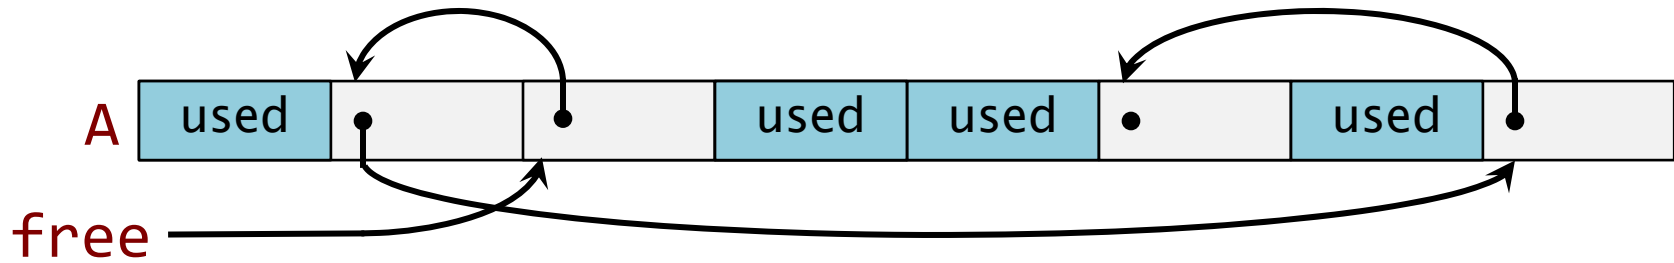
```
x = free;  
free = free->next;  
return x;
```

## free object x

```
x->next = free;  
free = x;
```

# Free Lists


## Free list



- Allocating and freeing take  $\Theta(1)$  time.
- Good temporal locality.
- Poor spatial locality due to **external fragmentation** — blocks distributed across virtual memory — which can increase the size of the page table and cause **disk thrashing**.
- The **translation lookaside buffer (TLB)** can also be a problem.


# Mitigating External Fragmentation

- Keep a free list (or bitmap) per disk page.
- Allocate from the free list for the fullest page.
- Free a block of storage to the free list for the page on which the block resides.
- If a page becomes empty (only free-list items), the virtual-memory system can page it out without affecting program performance.
- 90–10 is better than 50–50:



Probability that 2 random accesses hit the same page  
=  $.9 \times .9 + .1 \times .1 = .82$  versus  $.5 \times .5 + .5 \times .5 = .5$

# VARIABLE-SIZE HEAP ALLOCATION

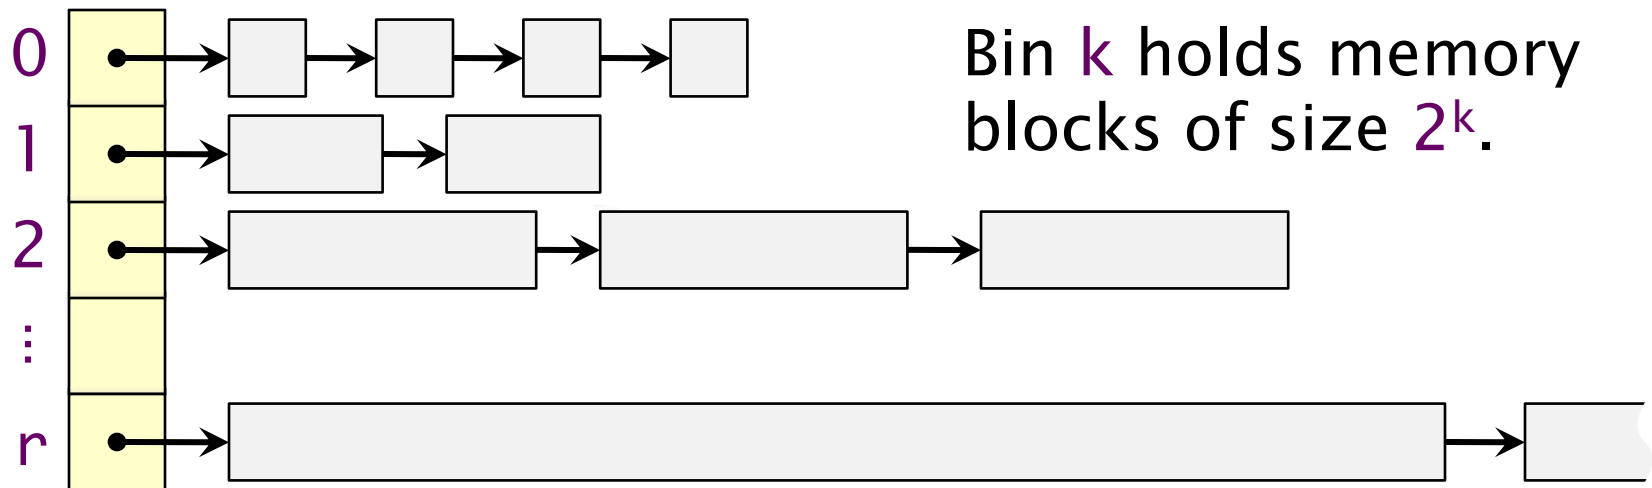




# Variable-Size Allocation

## Binned free lists


- Leverage the efficiency of free lists.
- Accept a bounded amount of internal fragmentation.



# Allocation for Binned Free Lists

Allocate  
 $x$  bytes

- If bin  $k = \lceil \lg x \rceil$  is nonempty, return a block.
- Otherwise, find a block in the next larger nonempty bin  $k' > k$ , split it up into blocks of sizes  $2^{k'-1}$ ,  $2^{k'-2}$ , ...,  $2^k$ ,  $2^k$ , and distribute the pieces.



Example


$x = 3 \Rightarrow \lceil \lg x \rceil = 2.$

Bin 2 is empty.

# Allocation for Binned Free Lists

Allocate  
 $x$  bytes

- If bin  $k = \lceil \lg x \rceil$  is nonempty, return a block.
- Otherwise, find a block in the next larger nonempty bin  $k' > k$ , split it up into blocks of sizes  $2^{k'-1}$ ,  $2^{k'-2}$ , ...,  $2^k$ ,  $2^k$ , and distribute the pieces.



Example


$x = 3 \Rightarrow \lceil \lg x \rceil = 2.$

Bin 2 is empty.

# Allocation for Binned Free Lists

Allocate  
 $x$  bytes

- If bin  $k = \lceil \lg x \rceil$  is nonempty, return a block.
- Otherwise, find a block in the next larger nonempty bin  $k' > k$ , split it up into blocks of sizes  $2^{k'-1}$ ,  $2^{k'-2}$ , ...,  $2^k$ ,  $2^k$ , and distribute the pieces.



Example

$x = 3 \Rightarrow \lceil \lg x \rceil = 2$ .

Bin 2 is empty.

return

# Allocation for Binned Free Lists

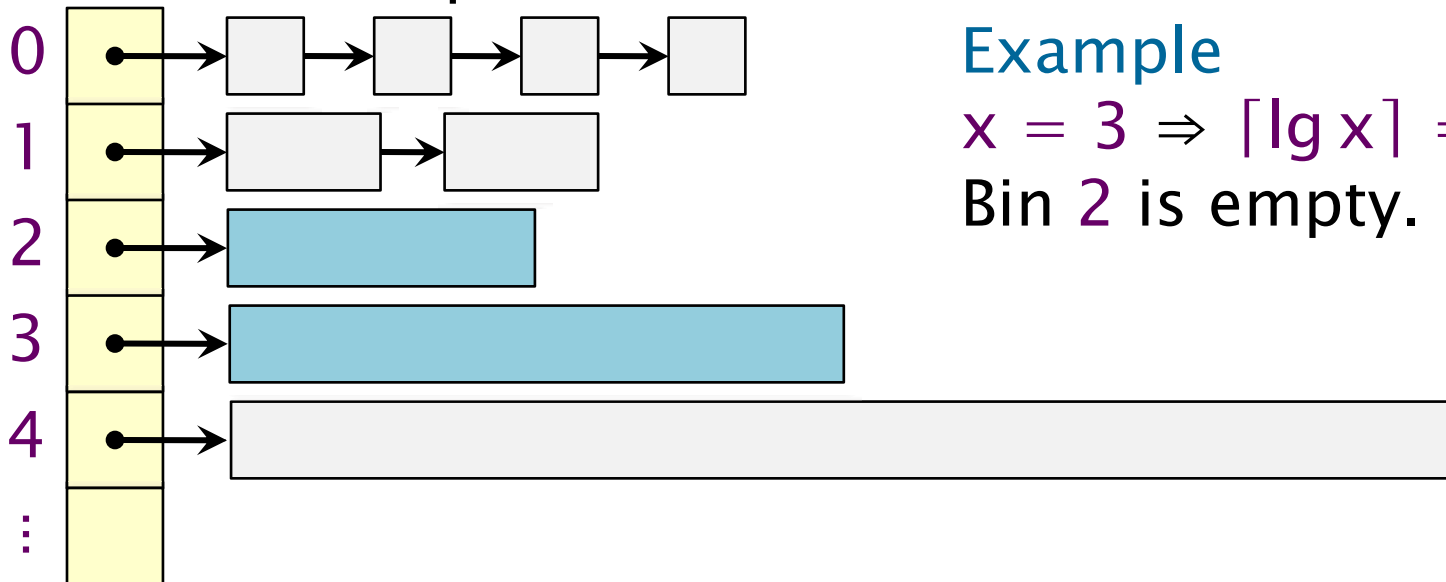
Allocate  
 $x$  bytes

- If bin  $k = \lceil \lg x \rceil$  is nonempty, return a block.
- Otherwise, find a block in the next larger nonempty bin  $k' > k$ , split it up into blocks of sizes  $2^{k'-1}$ ,  $2^{k'-2}$ , ...,  $2^k$ ,  $2^k$ , and distribute the pieces.\*

Example

$x = 3 \Rightarrow \lceil \lg x \rceil = 2$ .

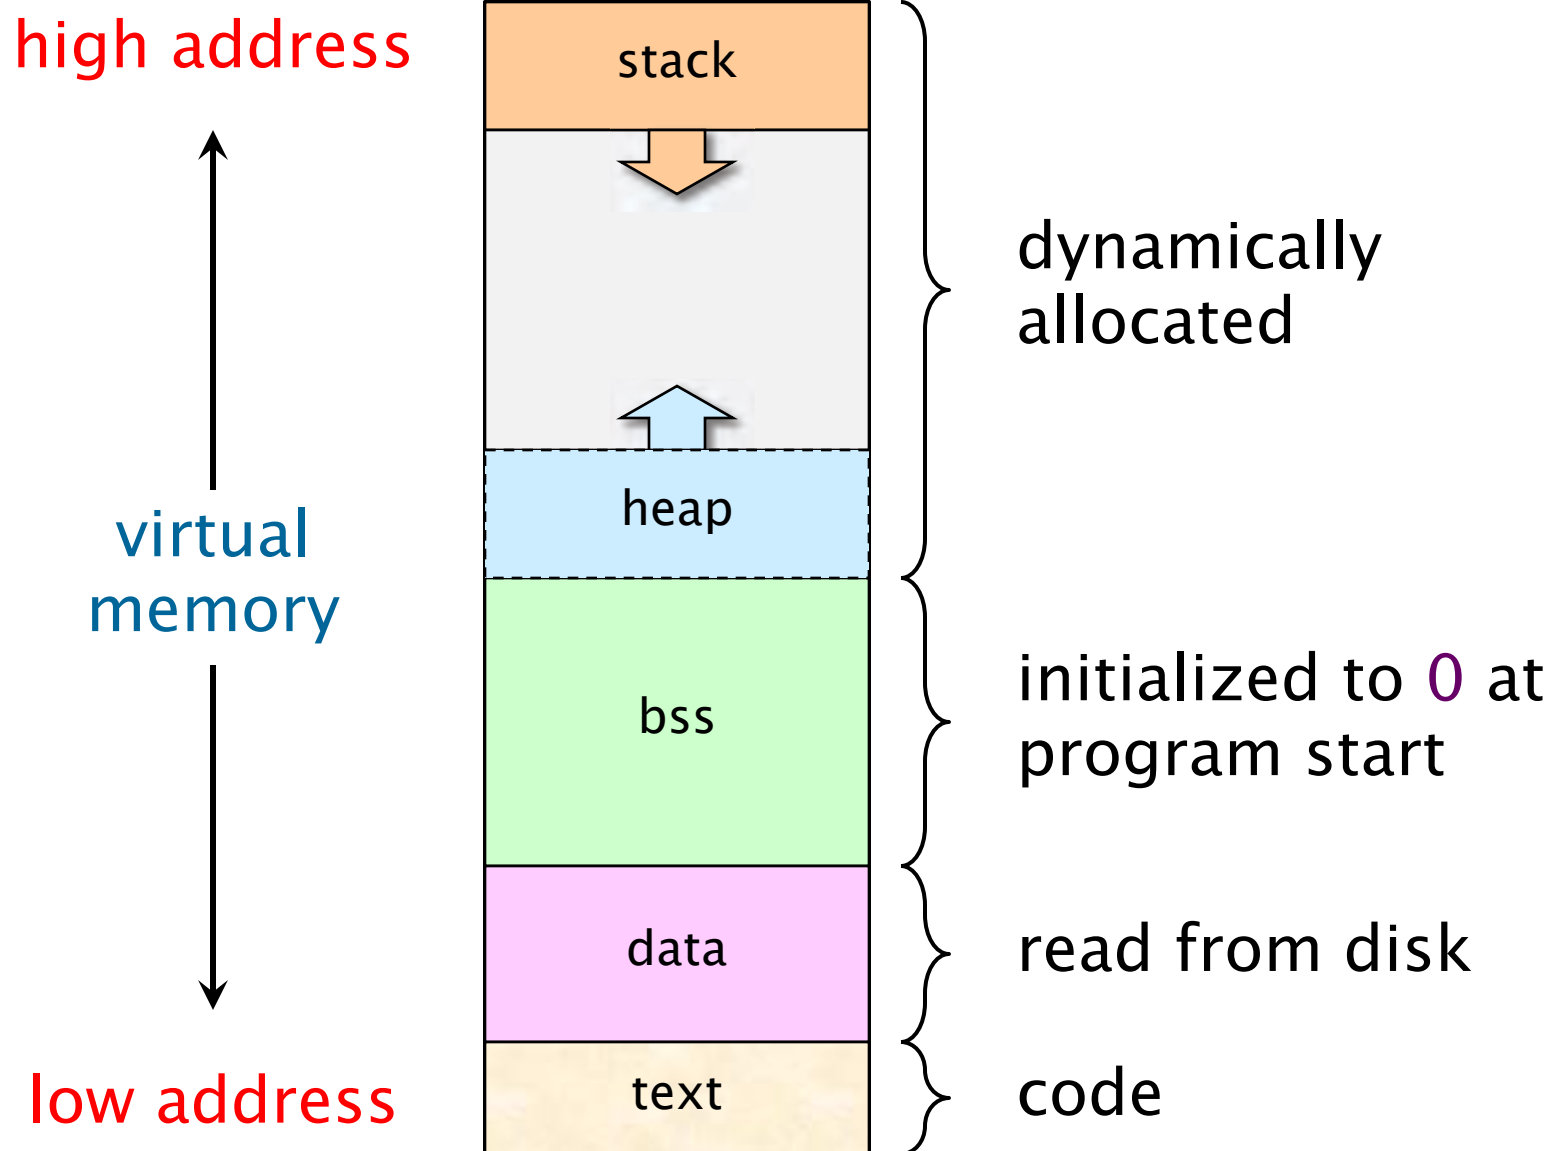
Bin 2 is empty.



return

\*If no larger blocks exist, ask the OS to allocate more memory.

# Storage Layout of a Program



# How Virtual is Virtual Memory?

- Q.** Since a 64-bit address space takes over a century to write at a rate of 4 billion bytes per second, we effectively never run out of virtual memory. Why not just allocate out of virtual memory and never free?
- A.** **External fragmentation** would be horrendous! The performance of the page table would degrade tremendously leading to **disk thrashing**, since all nonzero memory must be backed up on disk in page-sized blocks.

## Goal of storage allocators

Use as little virtual memory as possible, and try to keep the used portions relatively compact.

# Analysis of Binned Free Lists

**Theorem.** Suppose that the maximum amount of heap memory in use at any time by a program is  $M$ . If the heap is managed by a BFL allocator, the amount of virtual memory consumed by heap storage is  $O(M \lg M)$ .

*Proof.* An allocation request for a block of size  $x$  consumes  $2^{\lceil \lg x \rceil} \leq 2x$  storage. Thus, the amount of virtual memory devoted to blocks of size  $2^k$  is at most  $2M$ . Since there are at most  $\lg M$  free lists, the theorem holds. ■

⇒ In fact, BFL is  $\Theta(1)$ -competitive with the optimal allocator (assuming no coalescing).




# Coalescing

Binned free lists can sometimes be heuristically improved by **splicing together** adjacent small blocks into a larger block.

- Clever schemes exist for finding adjacent blocks efficiently — e.g., the “**buddy**” **system** — but the overhead is still greater than simple BFL.
- No good theoretical bounds exist that **prove** the effectiveness of coalescing.
- Coalescing seems to reduce fragmentation **in practice**, because heap storage tends to be deallocated as a stack (LIFO) or in batches.

# GARBAGE COLLECTION BY REFERENCE COUNTING



# Garbage Collectors

## Idea

- Free the programmer from freeing objects.
- A garbage collector identifies and recycles the objects that the program can no longer access.
- GC can be built-in (Java, Python) or do-it-yourself.



# Garbage Collection

## Terminology


- **Roots** are objects directly accessible by the program (globals, stack, etc.).
- **Live** objects are reachable from the roots by following pointers.
- **Dead** objects are inaccessible and can be recycled.

## How can the GC identify pointers?

- Strong typing.
- Prohibit pointer arithmetic (which may slow down some programs).


# Reference Counting

Keep a count of the number of pointers referencing each object. If the count drops to 0, free the dead object.




# Reference Counting

Keep a count of the number of pointers referencing each object. If the count drops to 0, free the dead object.




# Reference Counting

Keep a count of the number of pointers referencing each object. If the count drops to 0, free the dead object.



# Reference Counting


Keep a count of the number of pointers referencing each object. If the count drops to 0, free the dead object.






# Reference Counting

Keep a count of the number of pointers referencing each object. If the count drops to 0, free the dead object.



# Reference Counting


Keep a count of the number of pointers referencing each object. If the count drops to 0, free the dead object.



# Limitation of Reference Counting

## Problem

A cycle is never garbage collected!




# Limitation of Reference Counting

## Problem


A cycle is never garbage collected!

root





# MARK-AND-SWEEP GARBAGE COLLECTION




# Graph Abstraction

## Idea



Objects and pointers form a directed graph  $G = (V, E)$ . Live objects are reachable from the roots. Use breadth-first search to find the live objects.

## FIFO queue $Q$





```
for ( $\forall v \in V$ ) {  
    if (root( $v$ )) {  
         $v.mark = 1$ ;  
        enqueue( $Q, v$ );  
    } else  $v.mark = 0$ ;  
  
while ( $Q \neq \emptyset$ ) {  
     $u = \text{dequeue}(Q)$ ;  
    for ( $\forall v \in V$  such that  $(u, v) \in E$ ) {  
        if ( $v.mark == 0$ ) {  
             $v.mark = 1$ ;  
            enqueue( $Q, v$ );  
        }  
    }  
}
```

# Breadth-First Search









# Breadth-First Search





# Breadth-First Search





# Breadth-First Search





# Breadth-First Search





# Breadth-First Search





# Breadth-First Search



# Breadth-First Search

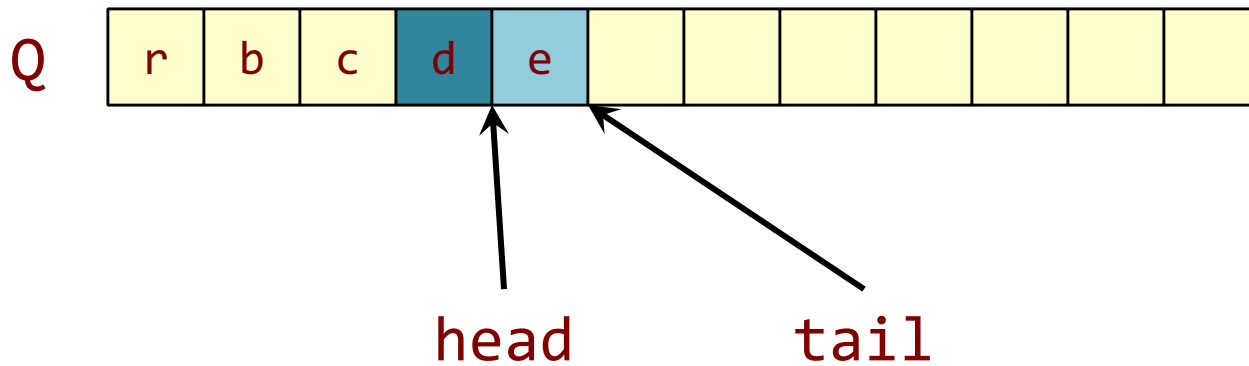



# Breadth-First Search









# Breadth-First Search





# Breadth-First Search





# Breadth-First Search





# Breadth-First Search





# Breadth-First Search



# Breadth-First Search



# Breadth-First Search



# Mark-and-Sweep


**Mark stage:** Breadth-first search marked all of the live objects.

**Sweep stage:** Scan over memory to free unmarked objects.



Mark-and-sweep doesn't deal with fragmentation



# STOP-AND-COPY GARBAGE COLLECTION





# Breadth-First Search




## Observation

All live vertices are placed in contiguous storage in  $Q$ .


# Copying Garbage Collector





# Copying Garbage Collector





# Copying Garbage Collector





# Copying Garbage Collector



# Copying Garbage Collector




# Copying Garbage Collector






# Copying Garbage Collector




When the **FROM** space is “full,” copy live storage using BFS with the **TO** space as the FIFO queue.

# Copying Garbage Collector



When the **FROM** space is “full,” copy live storage using BFS with the **TO** space as the FIFO queue.




# Updating Pointers

Since the **FROM** address of an object is not generally equal to the **TO** address of the object, pointers must be updated.


- When an object is copied to the **TO** space, store a forwarding pointer in the **FROM** object, which implicitly marks it as moved.
- When an object is removed from the FIFO queue in the **TO** space, update all its pointers.

# Example




Remove an item from the queue.

# Example




Remove an item from the queue.

# Example




Enqueue adjacent vertices.

# Example



Enqueue adjacent vertices.  
Place forwarding pointers in **FROM** vertices.


# Example



Update the pointers in the removed item to refer to its adjacent items in the **TO** space.




# Example




Update the pointers in the removed item to refer to its adjacent items in the **TO** space.

# Example



Linear time to copy and update all vertices.

# When Is the FROM Space “Full”?



- Request new heap space **equal to** the used space, and consider the FROM space to be “full” when this heap space has been allocated.
- The cost of garbage collection is then proportional to the size of the new heap space  $\Rightarrow$  amortized  $O(1)$  overhead, assuming that the user program touches all the memory allocated.
- Moreover, the VM space required is  $O(1)$  times optimal by locating the FROM and TO spaces in different regions of VM where they cannot interfere with each other.

# Dynamic Storage Allocation

Lots more is known and unknown about dynamic storage allocation. Strategies include

- buddy system,
- variants of mark-and-sweep,
- generational garbage collection,
- real-time garbage collection,
- multithreaded storage allocation,
- parallel garbage collection,
- etc.

# Summary

- Stack: most basic form of storage and is very efficient when it works
- Heap is the more general form of storage
- Fixed-size allocation using free lists
- Variable-sized allocation using binned free lists
- Garbage collection – reference counting, mark-and-sweep, stop-and-copy
- Internal and external fragmentation
- You will look at storage allocation in Homework 6 and Project 3