

6.172  
Performance  
Engineering  
of Software  
Systems



LECTURE 14

Caching and Cache-  
Efficient Algorithms

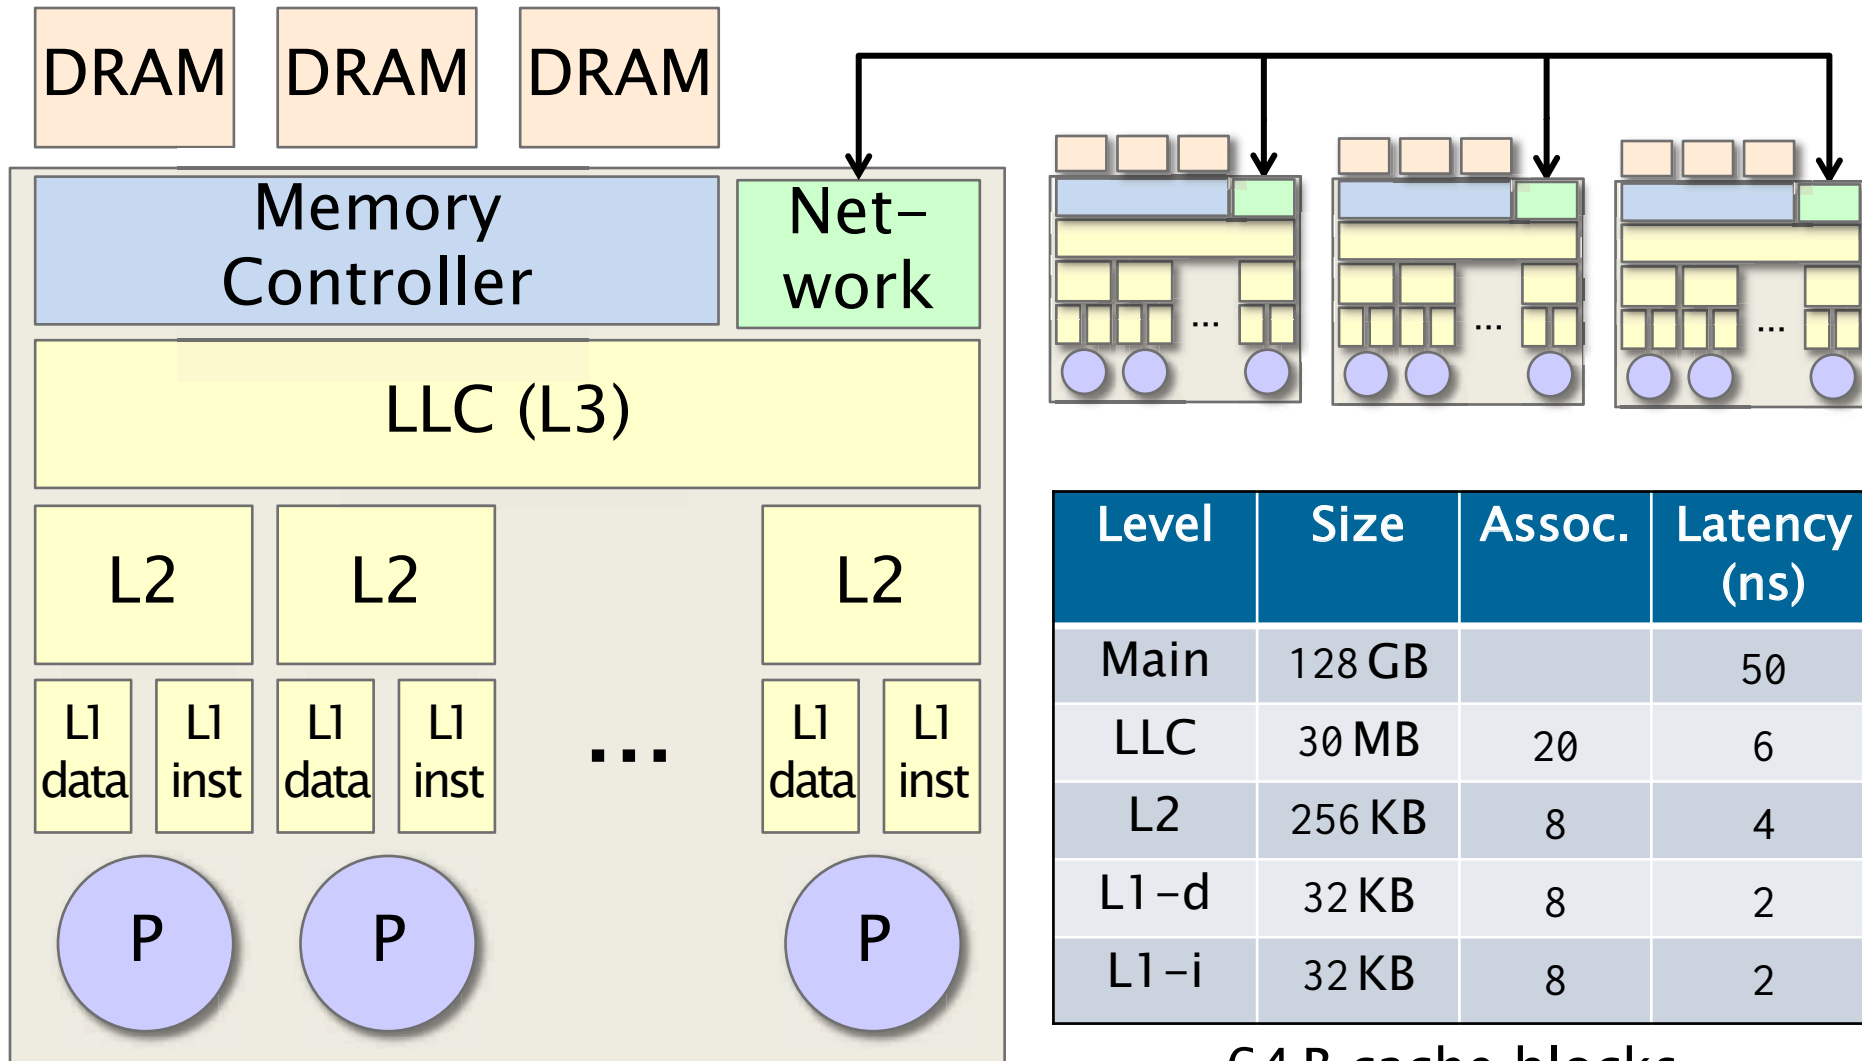
Julian Shun



# CACHE HARDWARE

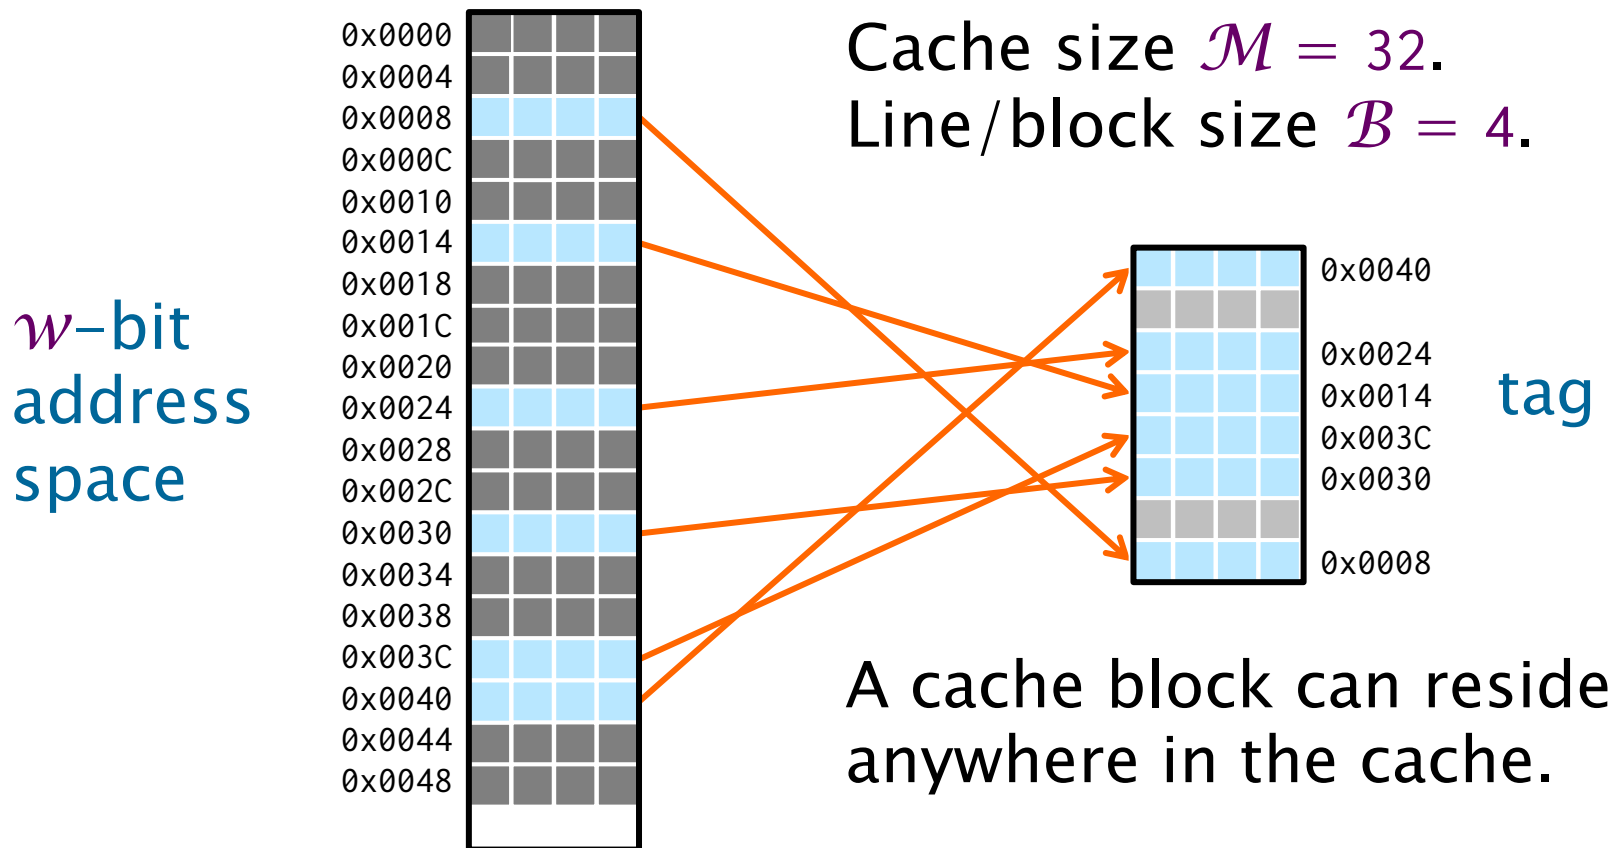


# Multicore Cache Hierarchy




64 B cache blocks

# Fully Associative Cache



To find a block in the cache, the entire cache must be searched for the tag. When the cache becomes full, a block must be **evicted** to make room for a new block. The **replacement policy** determines which block to evict.

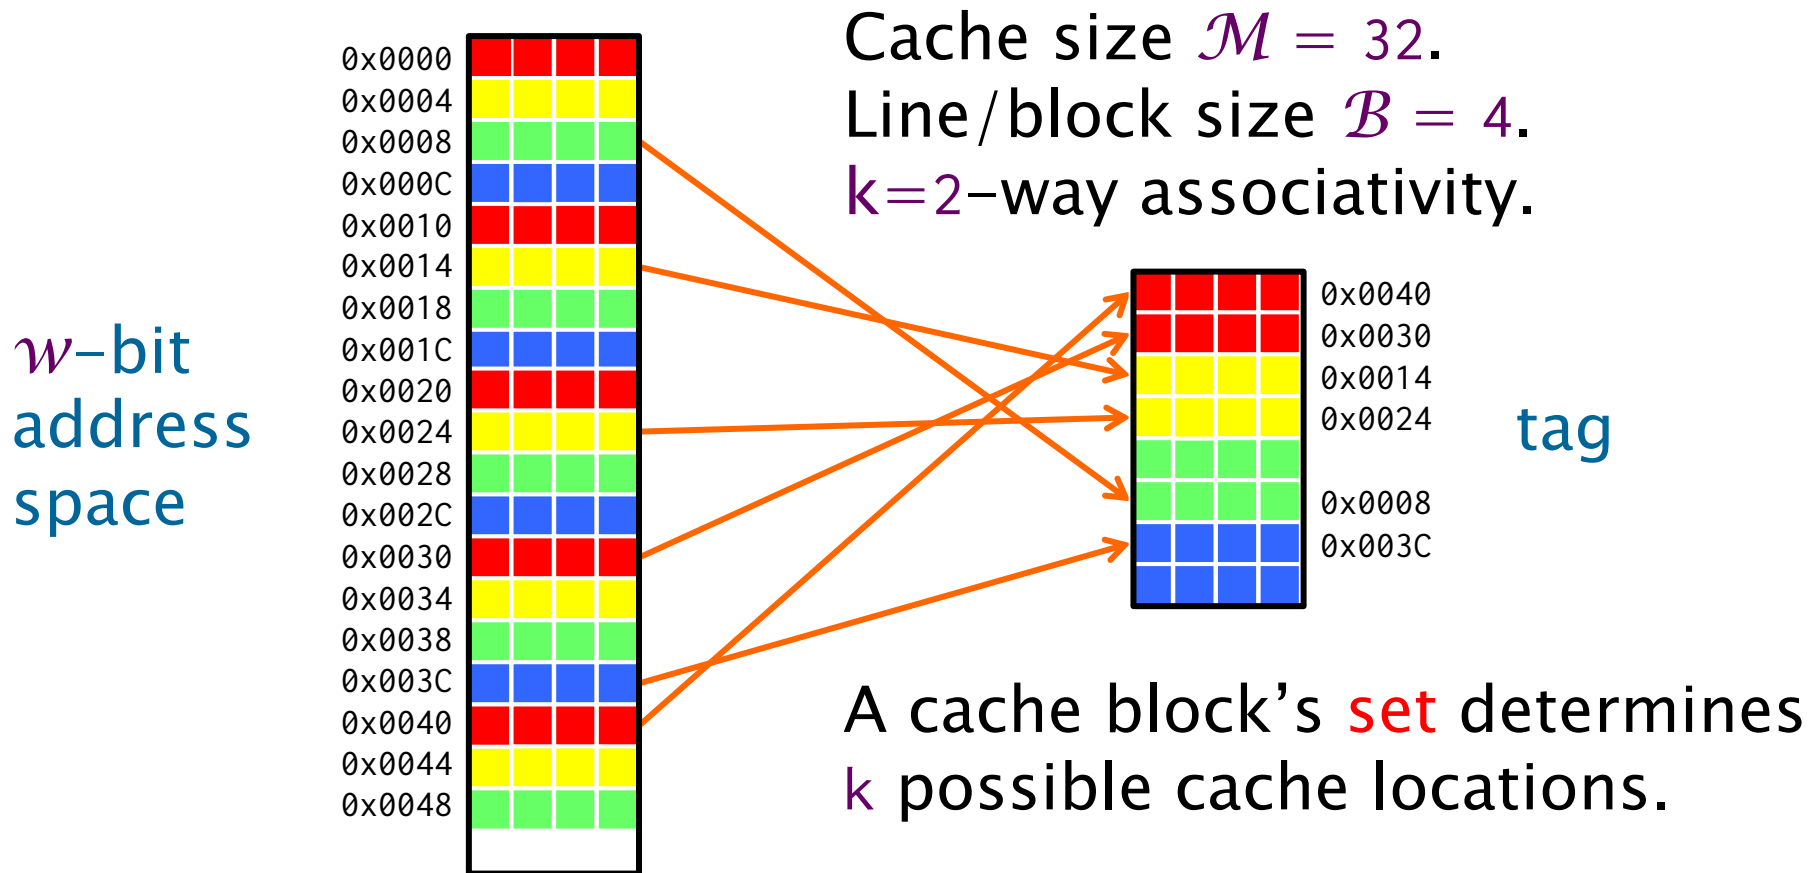
# Direct-Mapped Cache



address		
	tag	set
bits	$w - \lg \mathcal{M}$	$\lg(\mathcal{M}/\mathcal{B})$
	offset	
	$\lg \mathcal{B}$	

To find a block in the cache, only a single location in the cache need be searched.

# Set-Associative Cache



address		
	tag	set
bits	$w - \lg(\mathcal{M}/k)$	$\lg(\mathcal{M}/k\mathcal{B})$
	offset	
	$\lg \mathcal{B}$	

To find a block in the cache, only the  $k$  locations of its set must be searched.

# Taxonomy of Cache Misses

## Cold miss

- The first time the cache block is accessed.

## Capacity miss

- The previous cached copy would have been evicted even with a fully associative cache.


## Conflict miss

- Too many blocks from the same set in the cache. The block would not have been evicted with a fully associative cache.

## Sharing miss

- Another processor acquired exclusive access to the cache block.
- **True-sharing miss:** The two processors are accessing the same data on the cache line.
- **False-sharing miss:** The two processors are accessing different data that happen to reside on the same cache line.

# Conflict Misses for Submatrices



## Assume:

- Word width  $w = 64$ .
- Cache size  $\mathcal{M} = 32K$ .
- Line (block) size  $\mathcal{B} = 64$ .
- $k=4$ -way associativity.

Conflict misses can be problematic for caches with limited associativity.

## address

	tag	set	offset
bits	$w - \lg(\mathcal{M}/k)$	$\lg(\mathcal{M}/k\mathcal{B})$	$\lg \mathcal{B}$
	51	7	6

## Analysis

Look at a column of submatrix  $A$ . The addresses of the elements are  $x, x+2^{15}, x+2 \cdot 2^{15}, \dots, x+31 \cdot 2^{15}$ .

They all fall into the same set!

## Solutions

Copy  $A$  into a temporary  $32 \times 32$  matrix, or pad rows.




# IDEAL-CACHE MODEL



# Ideal-Cache Model

## Parameters

- Two-level hierarchy.
- Cache size of  $\mathcal{M}$  bytes.
- Cache-line length of  $\mathcal{B}$  bytes.
- Fully associative.
- Optimal, omniscient replacement.



## Performance Measures

- **work**  $\mathcal{W}$  (ordinary running time)
- **cache misses**  $\mathcal{Q}$

# How Reasonable Are Ideal Caches?

“**LRU**” **Lemma** [ST85]. Suppose that an algorithm incurs  $Q$  cache misses on an ideal cache of size  $\mathcal{M}$ . Then on a fully associative cache of size  $2\mathcal{M}$  that uses the **least-recently used (LRU)** replacement policy, it incurs at most  $2Q$  cache misses. ■

## Implication

For asymptotic analyses, one can assume optimal or LRU replacement, as convenient.

### Software Engineering

- Design a theoretically good algorithm.
- Engineer for detailed performance.
  - Real caches are not fully associative.
  - Loads and stores have different costs with respect to bandwidth and latency.

# Cache-Miss Lemma


**Lemma.** Suppose that a program reads a set of  $r$  data segments, where the  $i$ th segment consists of  $s_i$  bytes, and suppose that

$$\sum_{i=1}^r s_i = N < \mathcal{M}/3 \text{ and } N/r \geq \mathcal{B}.$$


Then all the segments fit into cache, and the number of misses to read them all is at most  $3N/\mathcal{B}$ .

**Proof.** A single segment  $s_i$  incurs at most  $s_i/\mathcal{B} + 2$  misses, and hence we have

$$\begin{aligned} \sum_{i=1}^r s_i/\mathcal{B} + 2 &= N/\mathcal{B} + 2r \\ &= (N/\mathcal{B} + 2r\mathcal{B})/\mathcal{B} \\ &\leq N/\mathcal{B} + 2N/\mathcal{B} \\ &= 3N/\mathcal{B}. \quad \blacksquare \end{aligned}$$



# Tall Caches




## Tall-cache assumption

$B^2 < c \mathcal{M}$  for some sufficiently small constant  $c \leq 1$ .

**Example:** Intel Xeon E5-2666 v3

- Cache-line length = 64 bytes.
- L1-cache size = 32 Kbytes.

# What's Wrong with Short Caches?




## Tall-cache assumption

$B^2 < cM$  for some sufficiently small constant  $c \leq 1$ .

An  $n \times n$  submatrix stored in row-major order may not fit in a short cache even if  $n^2 < cM$ !

# Submatrix Caching Lemma



**Lemma.** Suppose that an  $n \times n$  submatrix  $A$  is read into a tall cache satisfying  $\mathcal{B}^2 < c\mathcal{M}$ , where  $c \leq 1$  is constant, and suppose that  $c\mathcal{M} \leq n^2 < \mathcal{M}/3$ . Then  $A$  fits into cache, and the number of misses to read all  $A$ 's elements is at most  $3n^2/\mathcal{B}$ .

**Proof.** We have  $N = n^2$ ,  $n = r = s_i$ ,  $\mathcal{B} \leq n = N/r$ , and  $N < \mathcal{M}/3$ . Thus, the Cache-Miss Lemma applies. ■

# CACHE ANALYSIS OF MATRIX MULTIPLICATION





# Multiply Square Matrices

```
void Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i=0; i < n; i++)  
        for (int64_t j=0; j < n; j++)  
            for (int64_t k=0; k < n; k++)  
                C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```


## Analysis of work

$$W(n) = \Theta(n^3).$$


# Analysis of Cache Misses

```
void Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i=0; i < n; i++)  
        for (int64_t j=0; j < n; j++)  
            for (int64_t k=0; k < n; k++)  
                C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```

Assume row major and tall cache



A



B

Case 1

$n > c\mathcal{M}/\mathcal{B}$ .

Analyze matrix B.


Assume LRU.

$Q(n) = \Theta(n^3)$ , since matrix B misses on every access.


# Analysis of Cache Misses

```
void Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i=0; i < n; i++)  
        for (int64_t j=0; j < n; j++)  
            for (int64_t k=0; k < n; k++)  
                C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```

Assume row major and tall cache



A



B

## Case 2

$c' \mathcal{M}^{1/2} < n < c\mathcal{M}/\mathcal{B}$ .

Analyze matrix B.


Assume LRU.

$Q(n) = n \cdot \Theta(n^2/\mathcal{B}) = \Theta(n^3/\mathcal{B})$ , since matrix B can exploit spatial locality.


# Analysis of Cache Misses

```
void Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i=0; i < n; i++)  
        for (int64_t j=0; j < n; j++)  
            for (int64_t k=0; k < n; k++)  
                C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```

Assume row major and tall cache



A



B

## Case 3

$n < c' \mathcal{M}^{1/2}$ .

Analyze matrix B.


Assume LRU.

$Q(n) = \Theta(n^2 / \mathcal{B})$ ,  
since everything  
fits in cache!


# Swapping Inner Loop Order

```
void Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i=0; i < n; i++)  
        for (int64_t k=0; k < n; k++)  
            for (int64_t j=0; j < n; j++)  
                C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```

Assume row major and tall cache



C



B

Analyze matrix **B**.  
Assume LRU.


$Q(n) = n \cdot \Theta(n^2 / \mathcal{B}) = \Theta(n^3 / \mathcal{B})$ , since matrix **B** can exploit spatial locality.

TILING



# Tiled Matrix Multiplication

```
void Tiled_Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i1=0; i1<n/s; i1+=s)  
        for (int64_t j1=0; j1<n/s; j1+=s)  
            for (int64_t k1=0; k1<n/s; k1+=s)  
                for (int64_t i=i1; i<i1+s && i<n; i++)  
                    for (int64_t j=j1; j<j1+s && j<n; j++)  
                        for (int64_t k=k1; k<k1+s && k<n; k++)  
                            C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```




## Analysis of work

- Work  $W(n) = \Theta((n/s)^3(s^3))$   
 $= \Theta(n^3)$ .

# Tiled Matrix Multiplication

```
void Tiled_Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i1=0; i1<n; i1+=s)  
        for (int64_t j1=0; j1<n; j1+=s)  
            for (int64_t k1=0; k1<n; k1+=s)  
                for (int64_t i=i1; i<i1+s && i<n; i++)  
                    for (int64_t j=j1; j<j1+s && j<n; j++)  
                        for (int64_t k=k1; k<k1+s && k<n; k++)  
                            C[i*n+j] += A[i*n+k] * B[k*n+j];  
}
```



## Analysis of cache misses


- Tune  $s$  so that the submatrices just fit into cache  $\Rightarrow s = \Theta(\mathcal{M}^{1/2})$ .
- Submatrix Caching Lemma implies  $\Theta(s^2/\mathcal{B})$  misses per submatrix.
- $Q(n) = \Theta((n/s)^3(s^2/\mathcal{B}))$   
 $= \Theta(n^3/(\mathcal{B}\mathcal{M}^{1/2}))$ . *Remember this!*
- Optimal [HK81].



# Tiled Matrix Multiplication

```
void Tiled_Mult(double *C, double *A, double *B, int64_t n) {  
    for (int64_t i1  
        for (int  
            for (i  
                for  
                    for  
                        for (i  
                            C[i*n+j] += A[  
}
```


*Voodoo!*



## Analysis of cache misses


- Tune  $s$  so that the submatrices just fit into cache  $\Rightarrow s = \Theta(\mathcal{M}^{1/2})$ .
- Submatrix Caching Lemma implies  $\Theta(s^2/\mathcal{B})$  misses per submatrix.
- $Q(n) = \Theta((n/s)^3(s^2/\mathcal{B}))$   
 $= \Theta(n^3/(\mathcal{B}\mathcal{M}^{1/2}))$ . *Remember this!*
- Optimal [HK81].

# Two-Level Cache




- Two “voodoo” tuning parameters  $s$  and  $t$ .
- Multidimensional tuning optimization cannot be done with binary search.

# Two-Level Cache



```
void Tiled_Mult2(double *C, double *A, double *B, int64_t n) {
    for (int64_t i2=0; i2<n; i2+=s)
        for (int64_t j2=0; j2<n; j2+=s)
            for (int64_t k2=0; k2<n; k2+=s)
                for (int64_t i1=i2; i1<i2+s && i1<n; i1+=t)
                    for (int64_t j1=j2; j1<j2+s && j1<n; j1+=t)
                        for (int64_t k1=k2; k1<k2+s && k1<n; k1+=t)
                            for (int64_t i=i1; i<i1+s && i<i2+t && i<n; i++)
                                for (int64_t j=j1; j<j1+s && j<j2+t && j<n; j++)
                                    for (int64_t k=k1; k<k1+s && k<k2+t && k<n; k++)
                                        C[i*n+j] += A[i*n+k] * B[k*n+j];
}
```

# Three-Level Cache



- Three “voodoo” tuning parameters.
- Twelve nested for loops.
- Multiprogrammed environment: Don’t know the effective cache size when other jobs are running  $\Rightarrow$  easy to mistune the parameters!

# DIVIDE & CONQUER



# Recursive Matrix Multiplication

Divide-and-conquer on  $n \times n$  matrices.

$$\begin{array}{|c|c|} \hline C_{11} & C_{12} \\ \hline C_{21} & C_{22} \\ \hline \end{array} = \begin{array}{|c|c|} \hline A_{11} & A_{12} \\ \hline A_{21} & A_{22} \\ \hline \end{array} \times \begin{array}{|c|c|} \hline B_{11} & B_{12} \\ \hline B_{21} & B_{22} \\ \hline \end{array}$$
  
$$= \begin{array}{|c|c|} \hline A_{11}B_{11} & A_{11}B_{12} \\ \hline A_{21}B_{11} & A_{21}B_{12} \\ \hline \end{array} + \begin{array}{|c|c|} \hline A_{12}B_{21} & A_{12}B_{22} \\ \hline A_{22}B_{21} & A_{22}B_{22} \\ \hline \end{array}$$

8 multiply-adds of  $(n/2) \times (n/2)$  matrices.

# Recursive Code

```
// Assume that n is an exact power of 2.
void Rec_Mult(double *C, double *A, double *B,
              int64_t n, int64_t rowsize) {
    if (n == 1)
        C[0] += A[0] * B[0];
    else {
        int64_t d11 = 0;
        int64_t d12 = n/2;
        int64_t d21 = (n/2) * rowsize;
        int64_t d22 = (n/2) * (rowsize+1);


        Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
    }
}
```

Coarsen base case to overcome function-call overheads.

# Recursive Code

```
// Assume that n is an exact power of 2.
void Rec_Mult(double *C, double *A, double *B,
              int64_t n, int64_t rowsize) {
    if (n == 1)
        C[0] += A[0] * B[0];
    else {
        int64_t d11 = 0;
        int64_t d12 = n/2;
        int64_t d21 = (n/2) * rowsize;
        int64_t d22 = (n/2) * (rowsize+1);

        Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
    } }
```





# Analysis of Work

```
// Assume that n is an exact power of 2.
void Rec_Mult(double *C, double *A, double *B,
              int64_t n, int64_t rowsize) {
    if (n == 1)
        C[0] += A[0] * B[0];
    else {
        int64_t d11 = 0;
        int64_t d12 = n/2;
        int64_t d21 = (n/2) * rowsize;
        int64_t d22 = (n/2) * (rowsize+1);

        Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
    }
}
```

$$\begin{aligned} W(n) &= 8W(n/2) + \Theta(1) \\ &= \Theta(n^3) \end{aligned}$$

# Analysis of Work

$$W(n) = 8W(n/2) + \Theta(1)$$


recursion tree

$W(n)$

# Analysis of Work

$$W(n) = 8W(n/2) + \Theta(1)$$


recursion tree



# Analysis of Work


$$W(n) = 8W(n/2) + \Theta(1)$$

recursion tree



# Analysis of Work

$$W(n) = 8W(n/2) + \Theta(1)$$



**Note:** Same work as looping versions.

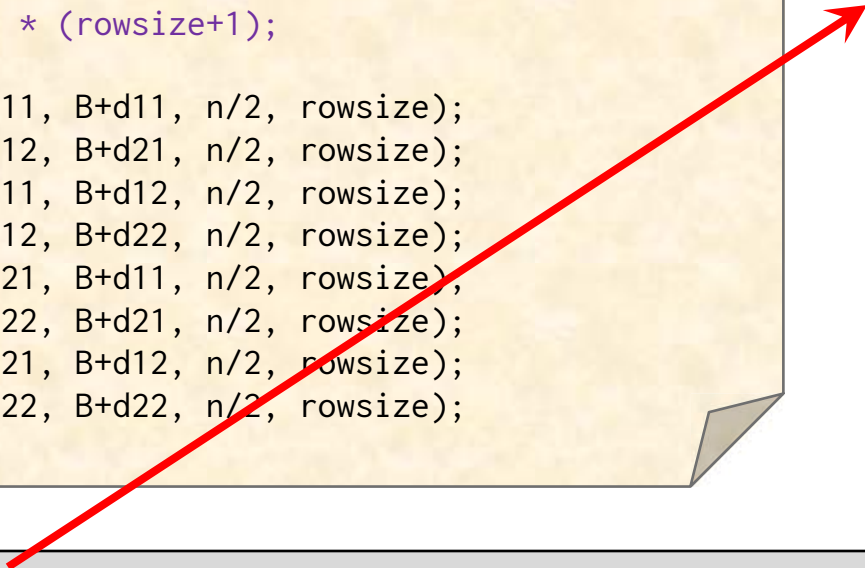
$$W(n) = \Theta(n^3)$$

# Analysis of Cache Misses

```
// Assume that n is an exact power of 2.
void Rec_Mult(double *C, double *A, double *B,
              int64_t n, int64_t rowsize) {
    if (n == 1)
        C[0] += A[0] * B[0];
    else {
        int64_t d11 = 0;
        int64_t d12 = n/2;
        int64_t d21 = (n/2) * rowsize;
        int64_t d22 = (n/2) * (rowsize+1);

        Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
    }
}
```

Submatrix  
Caching  
Lemma



$$Q(n) = \begin{cases} \Theta(n^2/\mathcal{B}) & \text{if } n^2 < c\mathcal{M} \text{ for suff. small const } c \leq 1, \\ 8Q(n/2) + \Theta(1) & \text{otherwise.} \end{cases}$$

# Analysis of Cache Misses

$$Q(n) = \begin{cases} \Theta(n^2 / \mathcal{B}) & \text{if } n^2 < c\mathcal{M} \text{ for suff. small const } c \leq 1, \\ 8Q(n/2) + \Theta(1) & \text{otherwise.} \end{cases}$$


recursion tree

$Q(n)$

# Analysis of Cache Misses

$$Q(n) = \begin{cases} \Theta(n^2/\mathcal{B}) & \text{if } n^2 < c\mathcal{M} \text{ for suff. small const } c \leq 1, \\ 8Q(n/2) + \Theta(1) & \text{otherwise.} \end{cases}$$

recursion tree






# Analysis of Cache Misses


$$Q(n) = \begin{cases} \Theta(n^2/\mathcal{B}) & \text{if } n^2 < c\mathcal{M} \text{ for suff. small const } c \leq 1, \\ 8Q(n/2) + \Theta(1) & \text{otherwise.} \end{cases}$$

recursion tree



# Analysis of Cache Misses

$$Q(n) = \begin{cases} \Theta(n^2/\mathcal{B}) & \text{if } n^2 < c\mathcal{M} \text{ for suff. small const } c \leq 1, \\ 8Q(n/2) + \Theta(1) & \text{otherwise.} \end{cases}$$



# Efficient Cache-Oblivious Algorithms

- No voodoo tuning parameters.
- No explicit knowledge of caches.
- Passively autotune.
- Handle multilevel caches automatically.
- Good in multiprogrammed environments.

## Matrix multiplication

The best cache-oblivious codes to date work on arbitrary rectangular matrices and perform binary splitting (instead of 8-way) on the largest of  $i$ ,  $j$ , and  $k$ .

# Recursive Parallel Matrix Multiply

```
// Assume that n is an exact power of 2.
void Rec_Mult(double *C, double *A, double *B,
              int64_t n, int64_t rowsize) {
    if (n == 1)
        C[0] += A[0] * B[0];
    else {
        int64_t d11 = 0;
        int64_t d12 = n/2;
        int64_t d21 = (n/2) * rowsize;
        int64_t d22 = (n/2) * (rowsize+1);

        cilk_spawn Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
        cilk_sync;
        cilk_spawn Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        cilk_sync;
    }
}
```

# Cilk and Caching

**Theorem.** Let  $Q_p$  be the number of cache misses in a deterministic Cilk computation when run on  $P$  processors, each with a private cache of size  $\mathcal{M}$ , and let  $S_p$  be the number of successful steals during the computation. In the ideal-cache model, we have

$$Q_p = Q_1 + O(S_p \mathcal{M} / \mathcal{B}) ,$$

where  $\mathcal{M}$  is the cache size and  $\mathcal{B}$  is the size of a cache block.

*Proof.* After a worker steals a continuation, its cache is completely cold in the worst case. But after  $\mathcal{M}/\mathcal{B}$  (cold) cache misses, its cache is identical to that in the serial execution. The same is true when a worker resumes a stolen subcomputation after a `cilk_sync`. The number of times these two situations can occur is at most  $2S_p$ . ■

$S_p = O(PT_\infty)$  in expectation

**MORAL:** Minimizing cache misses in the serial elision essentially minimizes them in parallel executions.

# Recursive Parallel Matrix Multiply

```
// Assume that n is an exact power of 2.
void Rec_Mult(double *C, double *A, double *B,
              int64_t n, int64_t rowsize) {
    if (n == 1)
        C[0] += A[0] * B[0];
    else {
        int64_t d11 = 0;
        int64_t d12 = n/2;
        int64_t d21 = (n/2) * rowsize;
        int64_t d22 = (n/2) * (rowsize+1);

        cilk_spawn Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
        cilk_sync;
        cilk_spawn Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        cilk_spawn Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        cilk_sync;
    }
}
```

*Span:*  $T_{\infty}(n) = 2T_{\infty}(n/2) + \Theta(1)$   
 $= \Theta(n)$

*Cache misses:*  $Q_p = Q_1 + O(S_p \mathcal{M}/B)$   
 $= \Theta(n^3 / B \mathcal{M}^{1/2}) + O(Pn \mathcal{M}/B)$

# Summary

- Associativity in caches
- Ideal cache model
- Cache-aware algorithms
  - Tiled matrix multiplication
- Cache-oblivious algorithms
  - Divide-and-conquer matrix multiplication
- Cache efficiency analysis in Homework 8