

High performance in dynamic languages:



6.172 guest lecture

Prof. [Steven G. Johnson](#)
MIT Applied Mathematics, MIT Physics

Dynamic languages for interactive math...

The two-language approach:

High-level dynamic language
for productivity,

+ low-level language (C,
Fortran, Cython, ...) for
performance-critical code.

= Huge jump in complexity,
loss of generality.

Just vectorize your code?

= rely on mature **external libraries**,
operating on **large blocks of data**,
for performance-critical code

Good advice! But...

- **Someone** has to write those libraries.
- Eventually that person will be **you**.
 - **some problems** are impossible or just very awkward to vectorize.

A new programming language?

Jeff Bezanson Viral Shah Alan Edelman
[MIT]
Stefan Karpinski

[30+ developers with 100+ commits,
1000+ external packages, 4th JuliaCon in 2017]



[begun 2009, “0.1” in 2013, ~40k commits,
“0.6” release in June 2017,
1.0 release in August 2018]

As **high-level and interactive** as Matlab or Python+IPython,
as **general-purpose** as Python,
as productive for **technical** work as Matlab or Python+SciPy,
but as **fast as C**.

Generating Vandermonde matrices

given $x = [\alpha_1, \alpha_2, \dots]$, generate:

$$V = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

NumPy ([numpy.vander](#)): *[follow links]*

[Python code](#) ...wraps [C code](#)
... wraps [generated C code](#)

type-generic at high-level, but
low level limited to small set of types.

Writing fast code “in” Python or Matlab = [mining the standard library](#)
for pre-written functions (implemented in C or Fortran).

If the problem doesn’t “vectorize” into built-in functions,
if you have to write your [own inner loops](#) ... [sucks](#) for you.

Generating Vandermonde matrices

given $x = [\alpha_1, \alpha_2, \dots]$, generate:

$$V = \begin{bmatrix} 1 & \alpha_1 & \alpha_1^2 & \dots & \alpha_1^{n-1} \\ 1 & \alpha_2 & \alpha_2^2 & \dots & \alpha_2^{n-1} \\ 1 & \alpha_3 & \alpha_3^2 & \dots & \alpha_3^{n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & \alpha_m & \alpha_m^2 & \dots & \alpha_m^{n-1} \end{bmatrix}$$

NumPy (`numpy.vander`): *[follow links]*


[Python code](#) ...wraps [C code](#)

... wraps [generated C code](#)


type-generic at high-level, but
low level limited to small set of types.

Julia (type-generic code):

```
function vander(x, n=length(x))
    m = length(x)
    V = Array{eltype(x), m, n}
    for j = 1:m
        V[j,1] = one(x[j])
    end
    for i = 2:n
        for j = 1:m
            V[j,i] = x[j] * V[j,i-1]
        end
    end
    return V
end
```



Generating Vandermonde matrices



Special Functions in Julia

Special functions $s(x)$: classic case that cannot be vectorized well

... switch between various polynomials depending on x

Many of Julia's special functions come from the usual C/Fortran libraries, but some are written in pure Julia code.

Pure Julia `erfinv(x)` [= $\text{erf}^{-1}(x)$]

3–4× faster than Matlab's and 2–3× faster than SciPy's (Fortran Cephes).

Pure Julia `polygamma(m, z)` [= $(m+1)^{\text{th}}$ derivative of the $\ln \Gamma$ function]

~ 2× faster than SciPy's (C/Fortran) for real z

... and unlike SciPy's, *same code* supports complex argument z

Julia code can actually be faster than typical “optimized” C/Fortran code, by using techniques [metaprogramming/codegen generation] that are hard in a low-level language.

Why can Julia be fast?

First need to understand: [Why is Python slow?](#)

goto Jupyter/Julia notebooks from 18.S096.