

What advantages does a set-associative cache have over a direct-mapped cache of the same total capacity?

- A Fewer cold misses.
- B Fewer capacity misses.
- C Less false sharing.
- D Fewer conflict misses.
- E Lower latency.

Solution:

D.

This is just knowing your definitions - look at lecture 14.

Which phenomena could cause the following code snippet to use the L1-cache ineffectively on a 12-core machine?

```
#define NUM_CHUNKS 12

int sums[NUM_CHUNKS];

void compute_sum(int chunk_id) {
    sums[chunk_id] = 0;
    for (int i = 0; i < 5000; i++) {
        sums[chunk_id] += i;
    }
}

int main() {
    cilk_for (int i = 0; i < NUM_CHUNKS; i++) {
        compute_sum(i);
    }
    return 0;
}
```

- A True sharing.
- B False sharing.
- C Capacity misses.
- D Conflict misses.
- E TLB misses.

Solution:

B.

All elements of the sums array are contiguous in memory, so adjacent entries share the same cache line. This leads to false sharing.

Consider the following code for multiplying two matrices:

```
// Assume that n is an exact power of 2.

void Rec_Mult(double *C, double *A, double *B,
              int n, int rowsize) {
    if (n == 1) {
        C[0] += A[0] * B[0];
    } else {
        int d11 = 0;
        int d12 = n/2;
        int d21 = (n/2) * rowsize;
        int d22 = (n/2) * (rowsize+1);

        Rec_Mult(C+d11, A+d11, B+d11, n/2, rowsize);
        Rec_Mult(C+d11, A+d12, B+d21, n/2, rowsize);
        Rec_Mult(C+d12, A+d11, B+d12, n/2, rowsize);
        Rec_Mult(C+d12, A+d12, B+d22, n/2, rowsize);
        Rec_Mult(C+d21, A+d21, B+d11, n/2, rowsize);
        Rec_Mult(C+d21, A+d22, B+d21, n/2, rowsize);
        Rec_Mult(C+d22, A+d21, B+d12, n/2, rowsize);
        Rec_Mult(C+d22, A+d22, B+d22, n/2, rowsize);
    }
}
```

For the following questions, assume that we have a tall cache of size \mathcal{M} , filled with cache lines of size \mathcal{B} , and assume that $\mathcal{B}^2 \leq c\mathcal{M}$ for some sufficiently small constant $c \leq 1$.

1. What recurrence does the cache complexity (number of cache misses) $Q(n)$ satisfy?
2. What is the height of the recurrence tree for $Q(n)$?
3. What is the total number of cache misses that occur in the leaves of the tree?
4. What is the total number of cache misses for the algorithm?

Solution:

1. What recurrence does the cache complexity (number of cache misses) $Q(n)$ satisfy?

If $n^2 < c\mathcal{M}$, then by submatrix caching lemma, it's $\Theta(n^2/\mathcal{B})$

Otherwise, it's $8Q(n/2) + \Theta(1)$

2. What is the height of the recurrence tree for $Q(n)$?

$\lg n - 1/2 * \lg(c\mathcal{M})$

3. What is the total number of cache misses that occur in the leaves of the tree?

number leaves = $8^{\lg n - 1/2 * \lg(c\mathcal{M})}$

= $\Theta(n^3 / (\mathcal{M}^{3/2}))$

Each leaf is $\Theta(c\mathcal{M}/\mathcal{B})$

Total cache misses is $\Theta(n^3 / (\mathcal{B} * \mathcal{M}^{1/2}))$

4. What is the total number of cache misses for the algorithm?

Dominated by leaves, so $\Theta(n^3 / (\mathcal{B} * \mathcal{M}^{1/2}))$